

# ALGORITHMS FOR NP-HARD PROBLEMS

## Part I – Heuristic Search

### Lecture 1: Introduction to Heuristic Search

Neil Yorke-Smith

[n.yorke-smith@tudelft.nl](mailto:n.yorke-smith@tudelft.nl)

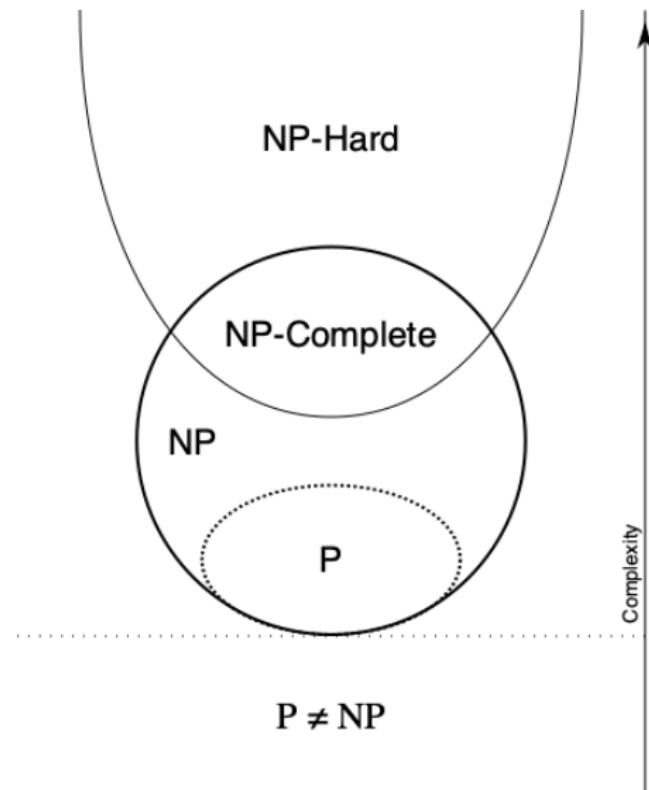


©2021–23 — CC BY-NC-NA 4.0

February 2023

# NP-hard

ACC



By Behnam Esfahbod, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=3532181>

# NP-hard

## Examples



Image by L.M. Ramirez-Elizondo,

<https://repository.tudelft.nl/islandora/object/uuid:3e2cb6d7-3ba2-4b45-af71-2fa106b5d189>

# NP-hard

## Examples



- making timetables (railways, schools, etc.)

Image by L.M. Ramirez-Elizondo,

<https://repository.tudelft.nl/islandora/object/uuid:3e2cb6d7-3ba2-4b45-af71-2fa106b5d189>

# NP-hard

## Examples



- making timetables (railways, schools, etc.)
- delivery of packages (e.g. TNT)

Image by L.M. Ramirez-Elizondo,

<https://repository.tudelft.nl/islandora/object/uuid:3e2cb6d7-3ba2-4b45-af71-2fa106b5d189>

# NP-hard

## Examples



- making timetables (railways, schools, etc.)
- delivery of packages (e.g. TNT)
- efficient routes for ride sharing (e.g. school bus, Valys)

Image by L.M. Ramirez-Elizondo,

<https://repository.tudelft.nl/islandora/object/uuid:3e2cb6d7-3ba2-4b45-af71-2fa106b5d189>

# NP-hard

## Examples



- making timetables (railways, schools, etc.)
- delivery of packages (e.g. TNT)
- efficient routes for ride sharing (e.g. school bus, Valys)
- scheduling processes (factories, hospitals, maintenance, CHPs in greenhouses, HVAC, etc.)



Image by L.M. Ramirez-Elizondo,

<https://repository.tudelft.nl/islandora/object/uuid:3e2cb6d7-3ba2-4b45-af71-2fa106b5d189>

# NP-hard

## Examples



- making timetables (railways, schools, etc.)
- delivery of packages (e.g. TNT)
- efficient routes for ride sharing (e.g. school bus, Valys)
- scheduling processes (factories, hospitals, maintenance, CHPs in greenhouses, HVAC, etc.)
- location of facilities (warehouses, shops, etc.)



Image by L.M. Ramirez-Elizondo,

<https://repository.tudelft.nl/islandora/object/uuid:3e2cb6d7-3ba2-4b45-af71-2fa106b5d189>

# Algorithms?



- What algorithms can we use?



# Algorithms?

- What algorithms can we use?
- Exponential time 😞

# Algorithms?



- What algorithms can we use?
- Exponential time 😞
- This course gives us new algorithmic tools



# Interactive block I: NP-hard



By Paddington\_station.location.map.svg: OpenStreetMap contributors (up by Edward)derivative work: Maly LOlek (talk) -  
Paddington\_station.location.map.svg, CC BY-SA 2.0, <https://commons.wikimedia.org/w/index.php?curid=6098976>



## Part I: Heuristic Search (Yorke-Smith)

How to get from A to B, probably as fast as you can.



## Part I: Heuristic Search (Yorke-Smith)

How to get from A to B, probably as fast as you can.

## Part II: Search and Inference (Demirović)

How to make exhaustive search less tiresome.



# About the course

## Part I: Heuristic Search (Yorke-Smith)

How to get from A to B, probably as fast as you can.

## Part II: Search and Inference (Demirović)

How to make exhaustive search less tiresome.

## Part III: Simulation-Based Planning (Lukina)

How to find approximate solutions under uncertainty.



# About the course

## Part I: Heuristic Search (Yorke-Smith)

How to get from A to B, probably as fast as you can.

## Part II: Search and Inference (Demirović)

How to make exhaustive search less tiresome.

## Part III: Simulation-Based Planning (Lukina)

How to find approximate solutions under uncertainty.

## Part IV (guest lecture): Dynamic Programming (Robu)

How to use what you've done before.



- Content on Brightspace



# Material, questions, assessment

- Content on Brightspace
- Questions on TUD Answers



- Content on Brightspace
- Questions on TUD Answers
- Assessment



- Content on Brightspace
- Questions on TUD Answers
- Assessment
  - Two optional lab sessions per part



# Material, questions, assessment

- Content on Brightspace
- Questions on TUD Answers
- Assessment
  - Two optional lab sessions per part
  - One assignment per part (in pairs)



# Material, questions, assessment

- Content on Brightspace
- Questions on TUD Answers
- Assessment
  - Two optional lab sessions per part
  - One assignment per part (in pairs)
  - One exam for all parts (individual, closed book)



- Content on Brightspace
- Questions on TUD Answers
- Assessment
  - Two optional lab sessions per part
  - One assignment per part (in pairs)
  - One exam for all parts (individual, closed book)
  - Grade =  $0.3 \cdot A + 0.7 \cdot E$



# Interactive block II: Questions



WHY DO WHALES JUMP  
WHY ARE WITCHES GREEN  
WHY ARE THERE MIRRORS ABOVE BEDS  
**WHY DO I SAY UH?**  
WHY IS SEA SALT BETTER  
WHY ARE THERE TREES IN THE MIDDLE OF FIELDS  
WHY IS THERE NOT A POKEMON MIMO  
WHY IS THERE LAUGHING IN TV SHOWS  
WHY ARE THERE DOORS ON THE FRENCH  
WHY ARE THERE SO MANY FRIENDS  
WHY AREN'T THERE ANY COUNTRIES IN ANTARCTICA  
WHY ARE THERE SCARY SOUNDS IN MINECRAFT  
WHY IS THERE KICKING IN MY STOMACH  
WHY ARE THERE TWO SUNSHIES AFTER HTTP  
WHY ARE THERE CELEBRITIES  
**WHY DO SNAKES EXIST?**  
WHY DO OYSTERS HAVE PEARLS  
WHY ARE DUCKS CALLED DUCKS  
WHY DO THEY CALL IT THE CLAP  
WHY ARE KYLE AND CARTMAN FRIENDS  
WHY IS THERE AN ARROW ON RANG'S HEAD  
WHY ARE TEXT MESSAGES BLUE  
WHY ARE THERE MUSTACHES ON CLOTHES  
WHY ARE THERE MUSTACHES ON CARS  
WHY ARE THERE MUSTACHES EVERYWHERE  
WHY ARE THERE SO MANY BIRDS IN OHIO  
WHY IS THERE SO MUCH RAIN IN OHIO  
WHY IS OHIO WEATHER SO WEIRD  
**WHY ARE THERE MALE AND FEMALE BIKES?**  
WHY ARE THERE BODYPOLIOS  
WHY DO DIVINE PEOPLE READ UP  
WHY ARE THERE UNICORN REINDEER  
WHY ARE OLD ALIENADS DIFFERENT  
**WHY ARE THERE SQUIRRELS?**  
WHY IS PROGRAMMING SO HIRED  
WHY IS THERE A 0 DM IN EDITOR  
WHY ARE THERE SO MANY SPIDERS  
WHY DO RHYMES SOUND GOOD  
**WHY DO TREES DIE?**  
WHY IS THERE NO EARTH DAY  
WHY ARE THERE SO MANY SPIDERS  
WHY AREN'T PERSONA PRACTICAL  
WHY AREN'T BULLETS SHARP  
WHY DO DREAMS SEEM SO REAL  
**WHY DO TESTICLES MOVE?**  
WHY ARE THERE PSYCHOS  
WHY ARE HATS SO EXPENSIVE  
WHY IS THERE COFFEE IN MY SHIRT  
WHY DO YOUR BOOBYS HURT?  
**WHY ARE THERE SLAVES IN THE BIBLE?**  
WHY DO TWINS HAVE DIFFERENT FINGERPRINTS  
WHY IS HTTPS CROSSED OUT IN RED  
WHY IS THERE A LINE THROUGH HTTPS  
WHY IS THERE A RED LINE THROUGH HTTPS ON FACEBOOK  
**WHY IS HTTPS IMPORTANT?**  
WHY AREN'T MY ARMS GROWING  
WHY ARE THERE SO MANY CROWS IN ROCHESTER, MN  
WHY ARE THERE SO MANY CROWS IN ROCHESTER, MN  
**WHY AREN'T ECONOMISTS RICH?**  
WHY DO AMERICANS CALL IT SOCCER  
**WHY ARE MY EARS RINGING?**  
WHY ARE THERE SO MANY AVENGERS  
WHY ARE THE AVENGERS FIGHTING THE X-MEN  
WHY IS WOLVERINE NOT IN THE AVENGERS  
**WHY ARE THERE ANTS IN MY LAPTOP?**  
WHY IS EARTH TILTED  
WHY IS SPACE BLACK  
WHY IS OUTER SPACE SO COLD  
WHY ARE THERE PYRAMIDS ON THE MOON  
WHY IS NASA SHUTTING DOWN  
WHY ARE THERE GHOSTS  
**WHY IS THERE AN OWL IN MY BACKYARD?**  
WHY IS THERE AN OWL OUTSIDE MY WINDOW  
WHY IS THERE AN OWL ON THE DOLLAR BILL  
**WHY DO OWLS ATTACK PEOPLE?**  
WHY ARE AK-47s SO EXPENSIVE  
WHY ARE THERE HELICOPTERS CIRCLING MY HOUSE  
WHY ARE THERE GODS  
WHY ARE THERE TWO SPOOKS  
**WHY IS MT VESUVIUS THERE?**  
WHY DO THEY SAY T MINUS  
WHY ARE THERE OBELISKS  
WHY ARE WRESTLERS ALWAYS WET  
WHY ARE OCEANS BECOMING MORE ACIDIC  
**WHY IS ARWEN DYING?**  
WHY AREN'T MY QUAIL LAYING EGGS  
WHY AREN'T MY QUAIL EGGS HATCHING  
WHY AREN'T THERE ANY FOREIGN MILITARY BASES IN AMERICA  
**WHY AREN'T THERE GUNS IN HARRY POTTER?**  
WHY AREN'T THERE GUNS IN HARRY POTTER  
WHY ARE ULTRASOUNDS IMPORTANT  
WHY ARE ULTRASOUNDS EXPENSIVE  
WHY IS STEALING WRONG

## QUESTIONS FOUND IN GOOGLE AUTOCOMPLETE

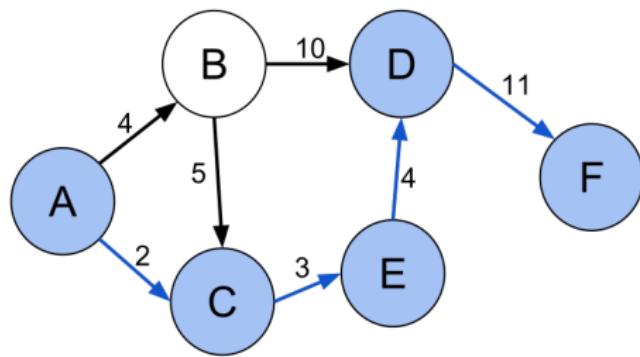
<https://xkcd.com/1256/>

# Why search?



<https://cs50.harvard.edu/ai/2020>

# Heuristic search



By Artyom Kalinin - Own work, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=29980338>

# Search trees

Search is at the heart of AI



<https://cs50.harvard.edu/ai/2020>

# **agent**

entity that perceives its environment  
and acts upon that environment

# **state**

a configuration of the agent and its environment

2	4	5	7
8	3	1	11
14	6		10
9	13	15	12

12	9	4	2
8	7	3	14
	1	6	11
5	13	10	15

15	4	10	3
13	1	11	12
9	5	14	7
6	8		2

# **initial state**

the state in which the agent begins

# initial state

2	4	5	7
8	3	1	11
14	6		10
9	13	15	12

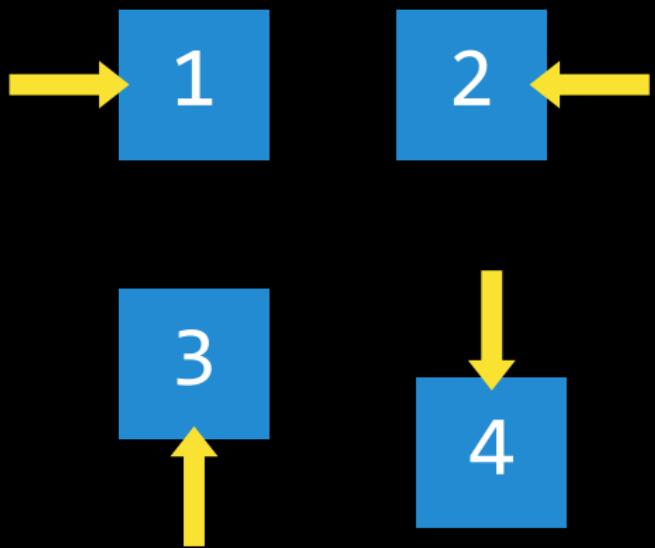
# **actions**

choices that can be made in a state

# actions

$\text{ACTIONS}(s)$  returns the set of actions that can be executed in state  $s$

actions



# **transition model**

a description of what state results from performing any applicable action in any state

# transition model

$\text{RESULT}(s, a)$  returns the state resulting from performing action  $a$  in state  $s$

RESULT(

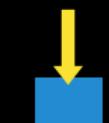
2	4	5	7
8	3	1	11
14	6	10	12
9	13	15	

,  ) =

2	4	5	7
8	3	1	11
14	6	10	12
9	13		15

RESULT(

2	4	5	7
8	3	1	11
14	6	10	12
9	13	15	

,  ) =

2	4	5	7
8	3	1	11
14	6	10	
9	13	15	12

# transition model

RESULT(

2	4	5	7
8	3	1	11
14	6	10	12
9	13	15	

,

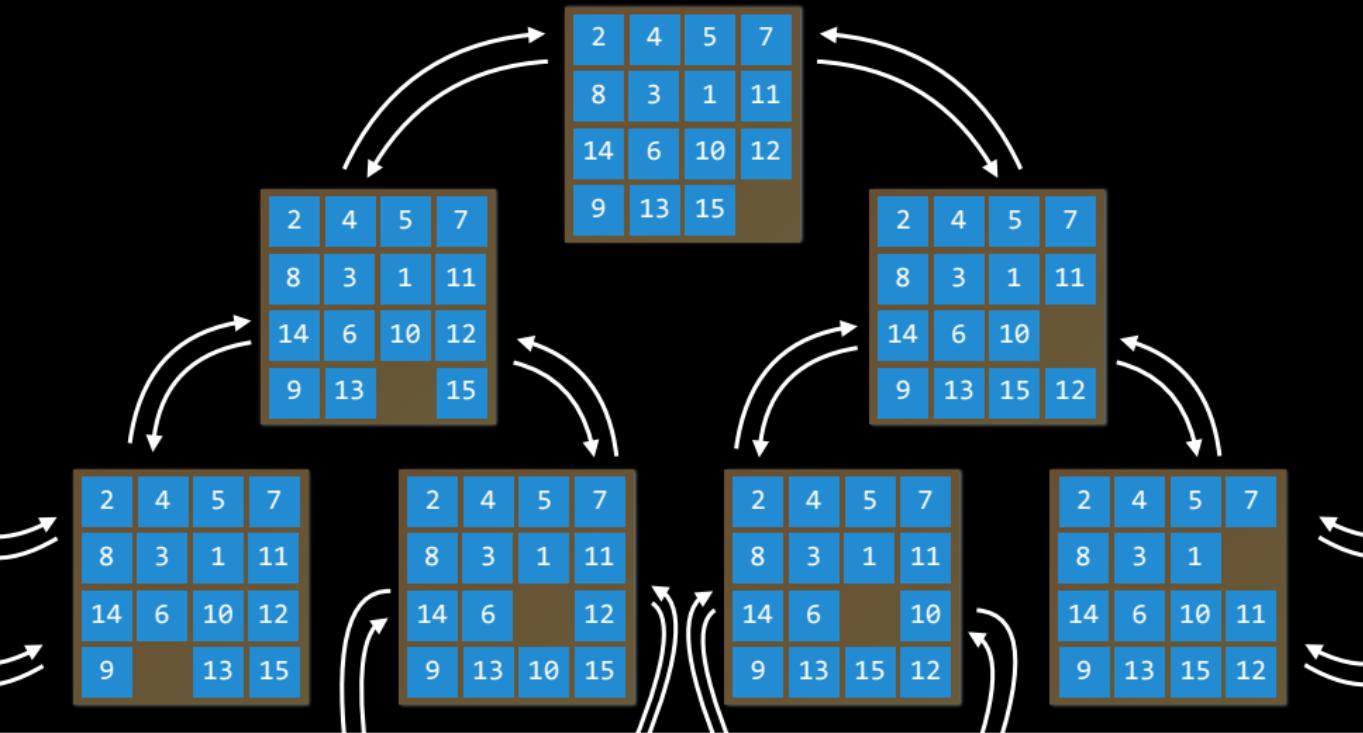


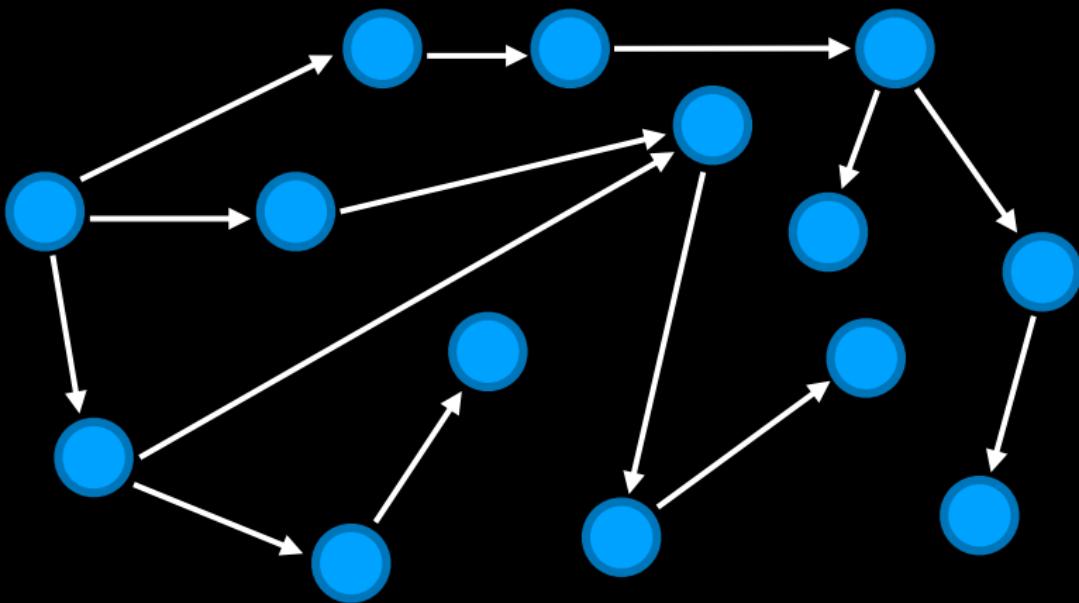
) =

2	4	5	7
8	3	1	11
14	6	10	
9	13	15	12

# **state space**

the set of all states reachable from the initial state by any sequence of actions



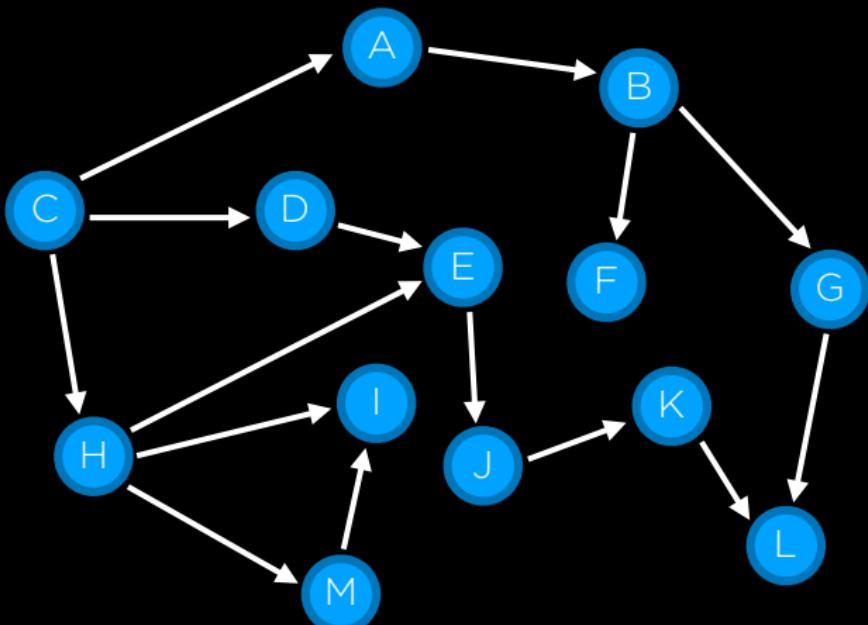


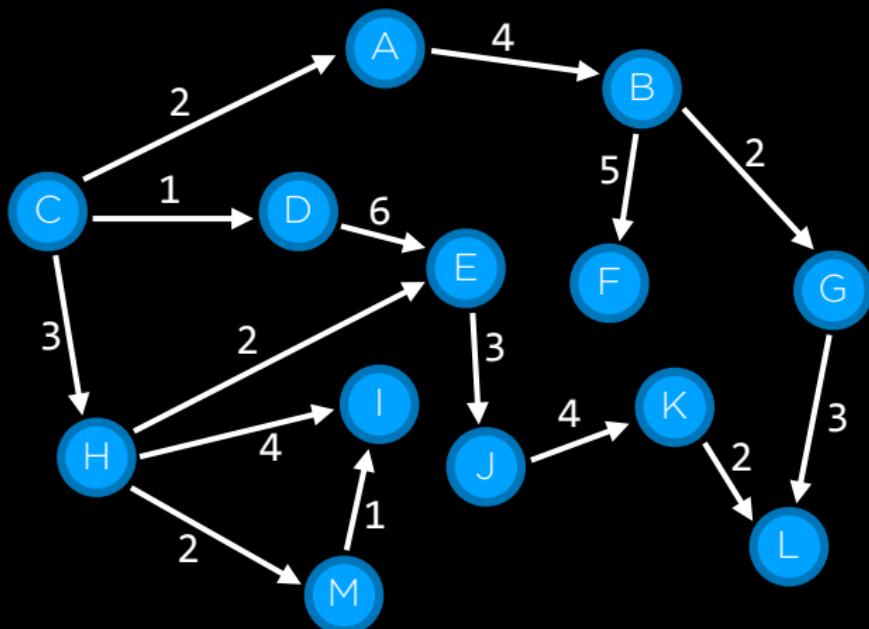
# **goal test**

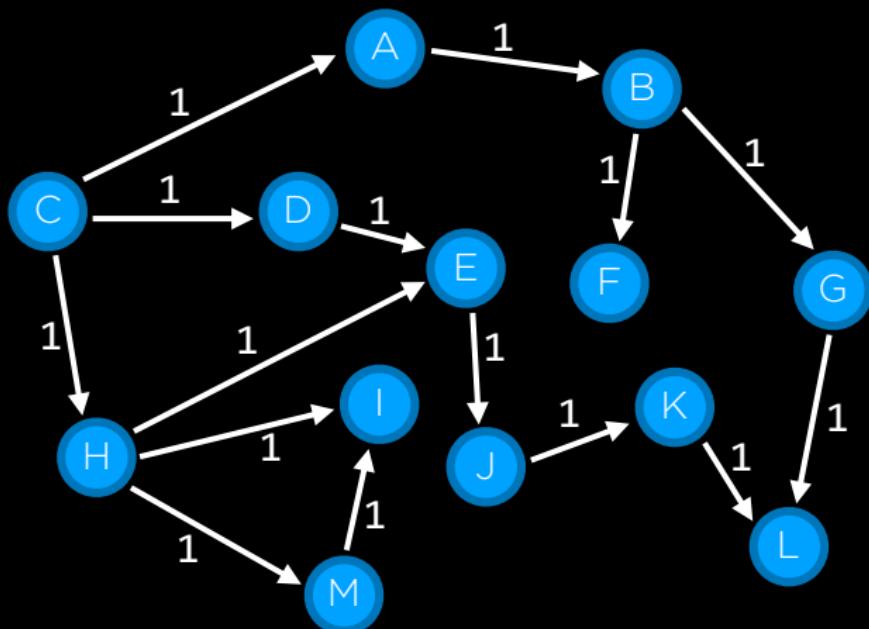
way to determine whether a given state  
is a goal state

# **path cost**

numerical cost associated with a given path







# Search Problems

- initial state
- actions
- transition model
- goal test
- path cost function

# Noughts-and-crosses

A game to analyse



By Symode09 - Own work, Public Domain, <https://commons.wikimedia.org/w/index.php?curid=2064271>



## Interactive block III: Search tree for OXO



Suppose  $X$  goes first, in the top-left corner

Write down the whole search tree for every possible move of  $O$  and  $X$ .



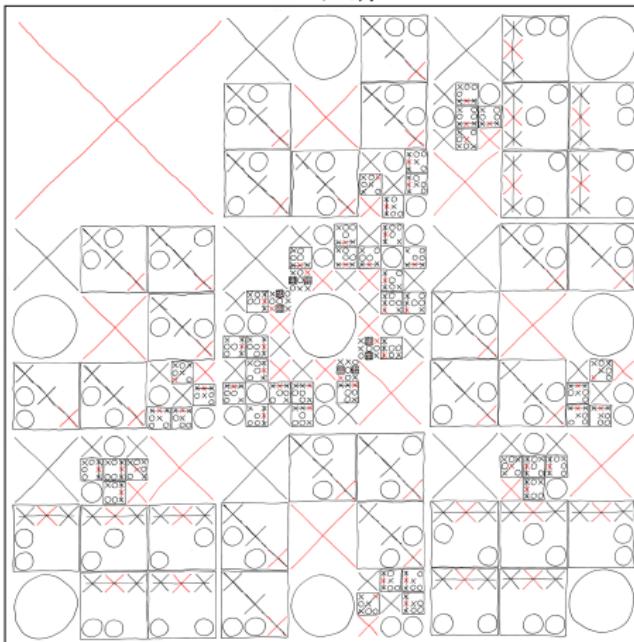
# Interactive block III: Search tree for OXO



## COMPLETE MAP OF OPTIMAL TIC-TAC-TOE MOVES

YOUR MOVE IS GIVEN BY THE POSITION OF THE LARGEST RED SYMBOL  
ON THE GRID. WHEN YOUR OPPONENT PICKS A MOVE, ZOOM IN ON  
THE REGION OF THE GRID WHERE THEY WENT. REPEAT.

MAP FOR X:



MAP FOR O:



<https://xkcd.com/832/>

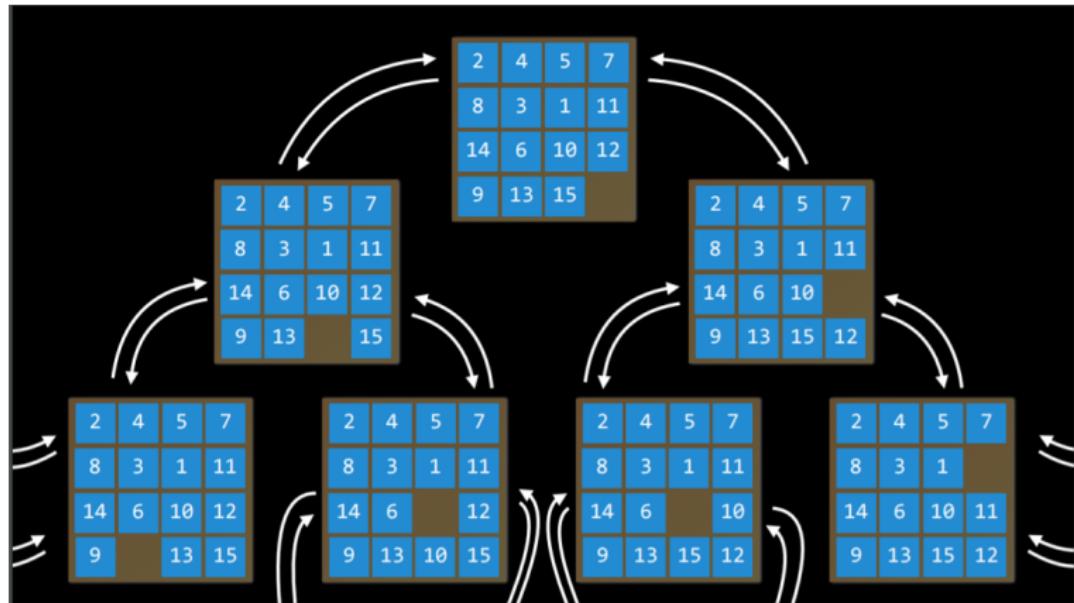


15 minute(s) break!



# Search trees

Search is at the heart of AI



<https://cs50.harvard.edu/ai/2020>

# Depth-First Search

Go deep



- Always explore a deepest node in the frontier

# Depth-First Search

Go deep



- Always explore a deepest node in the frontier
- Frontier = stack

# Depth-First Search

Go deep



- Always explore a deepest node in the frontier
- Frontier = stack
- Advantages?



## Interactive block IV: DFS for OXO

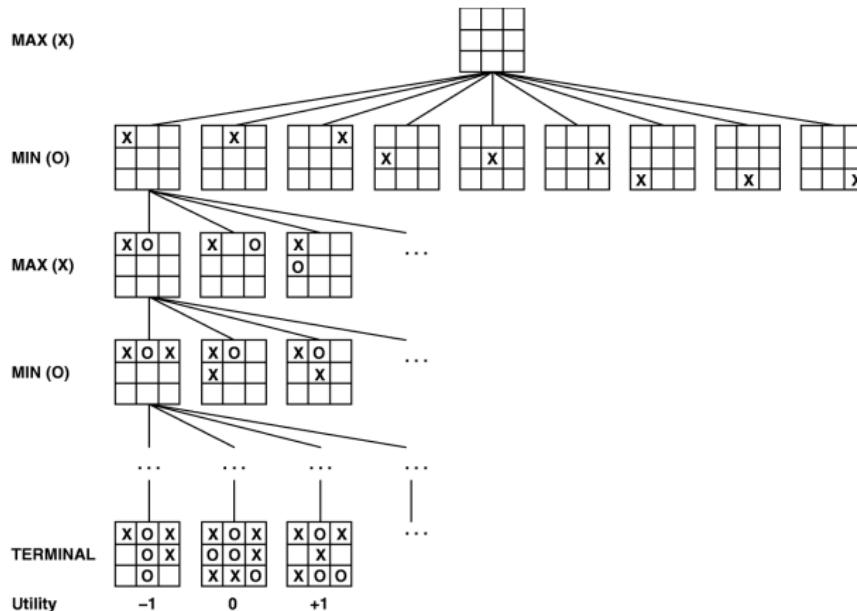


Suppose  $X$  goes first, in the top-left corner

Show a DFS traversal of your search tree.



## Interactive block IV: Search tree for OXO



<https://www.massey.ac.nz/~mjjohnso/>

# Breadth-First Search

Go wide



- Always explore a shallowest node in the frontier

# Breadth-First Search

Go wide



- Always explore a shallowest node in the frontier
- Frontier = queue

# Breadth-First Search

Go wide



- Always explore a shallowest node in the frontier
- Frontier = queue
- Advantages?

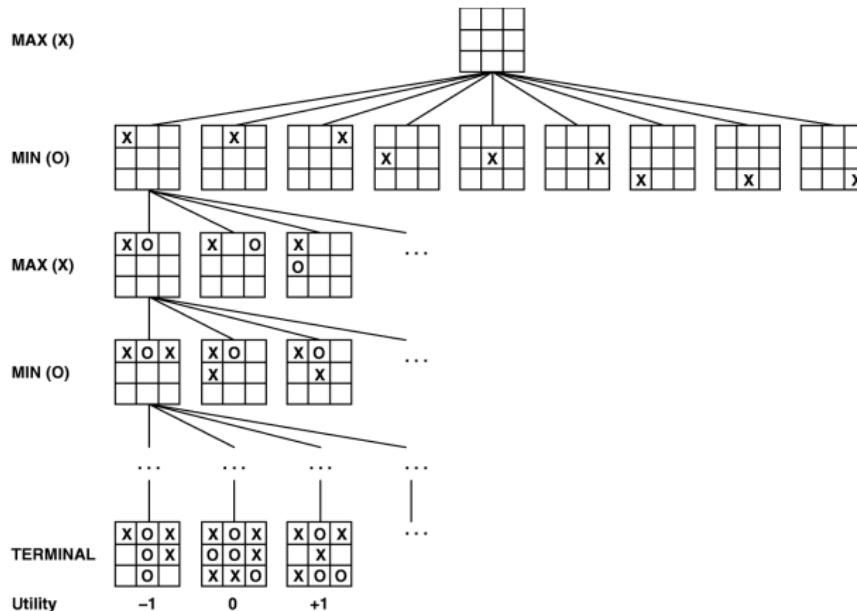


Suppose  $X$  goes first, in the top-left corner

Show a BFS traversal of your search tree.



# Interactive block V: Search tree for OXO



<https://www.massey.ac.nz/~mjjohnso/>

# Now what?



# Now what?

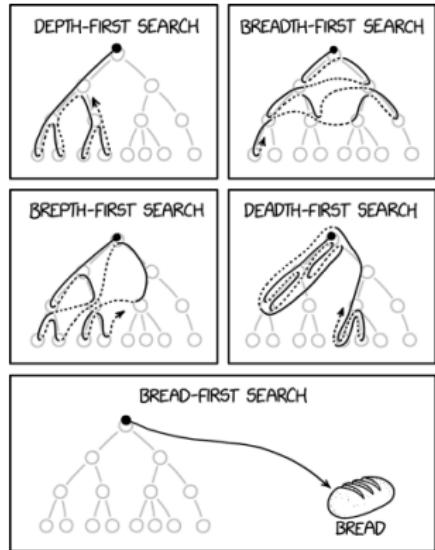


- DFS or BFS with smarter data structures?

# Now what?



- DFS or BFS with smarter data structures?
- Combining DFS and BFS?



<https://xkcd.com/2407/>

# Now what?



- DFS or BFS with smarter data structures?
- Combining DFS and BFS?
- DFS etc on graphs (instead of trees)?

# Now what?



- DFS or BFS with smarter data structures?
- Combining DFS and BFS?
- DFS etc on graphs (instead of trees)?
- DFS etc on hyper-graphs (instead of planar)?



# Now what?

- DFS or BFS with smarter data structures?
- Combining DFS and BFS?
- DFS etc on graphs (instead of trees)?
- DFS etc on hyper-graphs (instead of planar)?
- DFS etc in parallel?

# Now what?



- DFS or BFS with smarter data structures?
- Combining DFS and BFS?
- DFS etc on graphs (instead of trees)?
- DFS etc on hyper-graphs (instead of planar)?
- DFS etc in parallel?
- ...



## iBFS: Concurrent Breadth-First Search on GPUs

Hang Liu    H. Howie Huang    Yang Hu  
Department of Electrical and Computer Engineering  
George Washington University  
[{asherliu, howie, huyang}@gwu.edu](mailto:{asherliu, howie, huyang}@gwu.edu)

### ABSTRACT

Breadth-First Search (BFS) is a key graph algorithm with many important applications. In this work, we focus on a special class of graph traversal algorithm - concurrent BFS - where multiple breadth-first traversals are performed simultaneously on the same graph. We have designed and developed a new approach called *iBFS* that is able to run  $i$  concurrent BFSes from  $i$  distinct source vertices, very efficiently on Graphics Processing Units (GPUs). iBFS consists of three novel designs. First, iBFS develops a single GPU kernel for *joint traversal* of concurrent BFS to take advantage of shared frontiers across different instances. Second, *outdegree-based GroupBy rules* enables iBFS to selectively run a group of BFS instances which further maximizes the frontier sharing within such a group. Third, iBFS brings additional performance benefit by utilizing highly optimized *bitwise operations* on GPUs, which allows a single GPU thread to inspect a vertex for concurrent BFS instances. The evaluation on a wide spectrum of graph benchmarks shows that iBFS on one GPU runs up to 30x faster than executing BFS instances sequentially, and on 112 GPUs achieves near linear speedup with the maximum performance of 57,267 billion traversed edges per second (TEPS).

### 1. INTRODUCTION

Graph-based representations are widely applied in various fields such as social networks [1–3], metabolic networks [4],

Depending on the value of  $i$ , iBFS actually becomes a number of different problems. Formally, in a graph with  $|V|$  vertices, iBFS is

- single source shortest path (SSSP) if  $i = 1$  [6];
- multi-source shortest path (MSSP) if  $i \in (1, |V|)$  [7,8];
- all-pairs shortest path (APSP) if  $i = |V|$  [9,10].

Moreover, iBFS can be utilized in many other graph algorithms such as betweenness centrality [11,12] and closeness centrality [13]. For example, one can leverage iBFS to construct the index for answering graph reachability queries, that is, whether there exists a path from vertex  $s$  to  $t$  with the number of edges in-between less than  $k$  [14–16]. We will show in this paper this step can be an order of magnitude faster with iBFS. In all, a wide variety of applications, e.g., network routing [17–19], network attack detection [20], route planning [21–23], web crawling [24,25], etc., can benefit from high-performance iBFS.

This work proposes a new approach for running iBFS on GPUs that consists of three novel techniques: joint traversal, GroupBy, and bitwise optimization. Prior work has proposed to combine the execution of different BFS instances mostly on multi-core CPUs [26–28]. The performance improvement, however, is limited. For one, none of early projects has attempted to group the BFS instances to improve the frontier sharing during the traversal. Further, bottom-up BFS on the GPU has many challenges. For example, while MS-BFS [26] supports bottom-up, it does not provide early

# Now what?





## Informed search

What if we have some problem-specific information?

# (Greedy) Best-First Search

Go smart



- Always explore a best node in the frontier

# (Greedy) Best-First Search

Go smart



- Always explore a best node in the frontier
- Best = measured by an estimate  $h(n)$  of how close node  $n$  is to the goal

# (Greedy) Best-First Search

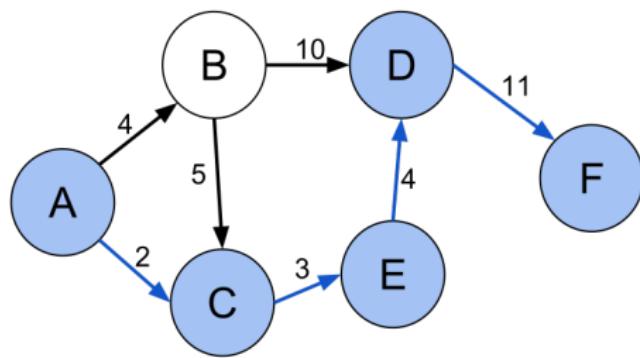
Go smart



- Always explore a best node in the frontier
- Best = measured by an estimate  $h(n)$  of how close node  $n$  is to the goal
- Advantages?

# Best node = closest to goal

According to heuristic  $h(n)$



By Artyom Kalinin - Own work, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=29980338>



Suppose  $X$  goes first, in the top-left corner

Use the heuristic function:

$$h(n) = \sum_{x \in \text{squares}} \begin{cases} 3 & x = \text{middle} \\ 2 & x = \text{corner} \\ 1 & x = \text{otherwise} \end{cases} - \sum_{o \in \text{squares}} \begin{cases} 3 & o = \text{middle} \\ 2 & o = \text{corner} \\ 1 & o = \text{otherwise} \end{cases}$$

Show a BestFS traversal of your search tree.



## Interactive block VI: BestFS for OXO

Example heuristic function



<https://cs50.harvard.edu/ai/2020>



## Interactive block VI: BestFS for OXO

Example heuristic function



$$h(n) =$$

<https://cs50.harvard.edu/ai/2020>



## Interactive block VI: BestFS for OXO

Example heuristic function



$$h(n) = 3 + 2 + 1 -$$

<https://cs50.harvard.edu/ai/2020>



## Interactive block VI: BestFS for OXO

Example heuristic function



$$h(n) = 3 + 2 + 1 - (2 + 2 + 1) =$$

<https://cs50.harvard.edu/ai/2020>



## Interactive block VI: BestFS for OXO

Example heuristic function



$$h(n) = 3 + 2 + 1 - (2 + 2 + 1) = 1$$

<https://cs50.harvard.edu/ai/2020>



## Interactive block VI: BestFS for OXO

Example heuristic function



$$h(n) = 3 + 2 + 1 - (2 + 2 + 1) = 1$$

What terminal states is this heuristic designed to find?

<https://cs50.harvard.edu/ai/2020>

## Next lecture



- Today we saw basic tree search



- Today we saw basic tree search
  - Heuristics for game playing Lectures 3 & 4



- Today we saw basic tree search
  - Heuristics for game playing ↗ Lectures 3 & 4
  - Exhaustive tree search will return in Part 2



# Next lecture

- Today we saw basic tree search
  - Heuristics for game playing  Lectures 3 & 4
  - Exhaustive tree search will return in Part 2
- Next time: smarter greedy best-first search



- Today we saw basic tree search
  - Heuristics for game playing Lectures 3 & 4
  - Exhaustive tree search will return in Part 2
- Next time: smarter greedy best-first search
- Preparation: see Brightspace



## Next lecture

- Today we saw basic tree search
  - Heuristics for game playing Lectures 3 & 4
  - Exhaustive tree search will return in Part 2
- Next time: smarter greedy best-first search
- Preparation: see Brightspace
- Lab on Monday



# Next lecture

- Today we saw basic tree search
  - Heuristics for game playing  Lectures 3 & 4
  - Exhaustive tree search will return in Part 2
- Next time: smarter greedy best-first search
- Preparation: see Brightspace
- Lab on Monday
- Questions? TUD Answers



# What is still unclear?

After every lecture...

Give us some homework:  
tell us what is still unclear.

# ALGORITHMS FOR NP-HARD PROBLEMS

## Part I – Heuristic Search

### Lecture 2: Greedy Best-First Search

Neil Yorke-Smith

[n.yorke-smith@tudelft.nl](mailto:n.yorke-smith@tudelft.nl)



©2021-23 — CC BY-NC-NA 4.0

February 2023



# Where are we?

- Last time: systematically searching a tree
- Today: (more) informed search



# Where are we?

- Last time: systematically searching a tree
- Today: (more) informed search
- You prepared:
  - Video on A\* search

# Questions from Lecture 1



# Uninformed search

Uses no problem-specific knowledge



<https://cs50.harvard.edu/ai/2020>

# Informed search

Uses problem-specific knowledge



Informed search

Heuristic function estimates

# Informed search

Uses problem-specific knowledge



## Informed search

Heuristic function estimates distance of node to goal.

# Informed search

Uses problem-specific knowledge



## Informed search

Heuristic function estimates distance of node to goal.

Is informed better than uninformed?

# Informed search

Uses problem-specific knowledge



## Informed search

Heuristic function estimates distance of node to goal.

Is informed better than uninformed?

Not always.

# Dijkstra's algorithm

aka Uniform Cost Search



By Hamilton Richards - manuscripts of Edsger W. Dijkstra, University Texas at Austin, CC BY-SA 3.0,  
<https://commons.wikimedia.org/w/index.php?curid=4204157>

# Dijkstra's algorithm

aka Uniform Cost Search



- Expands from starting node

By Hamilton Richards - manuscripts of Edsger W. Dijkstra, University Texas at Austin, CC BY-SA 3.0,  
<https://commons.wikimedia.org/w/index.php?curid=4204157>

# Dijkstra's algorithm

aka Uniform Cost Search



- Expands from starting node
- Uniformly, using cost-from-start-to-current-node  $g(n)$

By Hamilton Richards - manuscripts of Edsger W. Dijkstra, University Texas at Austin, CC BY-SA 3.0,  
<https://commons.wikimedia.org/w/index.php?curid=4204157>

# Dijkstra's algorithm

aka Uniform Cost Search



- Expands from starting node
- Uniformly, using cost-from-start-to-current-node  $g(n)$
- Guarantees shortest path (unlike Greedy BestFS)

By Hamilton Richards - manuscripts of Edsger W. Dijkstra, University Texas at Austin, CC BY-SA 3.0,  
<https://commons.wikimedia.org/w/index.php?curid=4204157>



- Look back and look forward (with estimate)



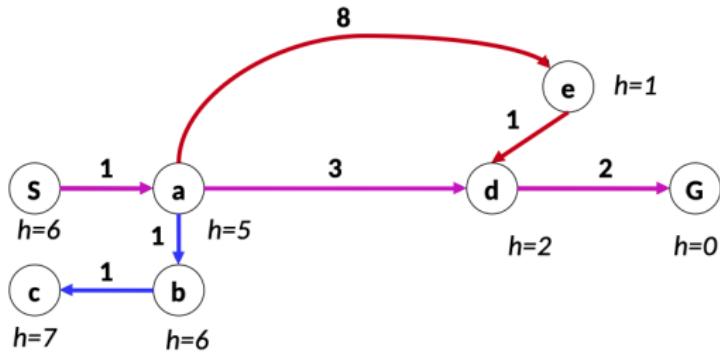
- Look back and look forward (with estimate)
- Use  $g(n) + h(n)$



- Look back and look forward (with estimate)
- Use  $g(n) + h(n)$
- Advantages?

# A\* search

Example

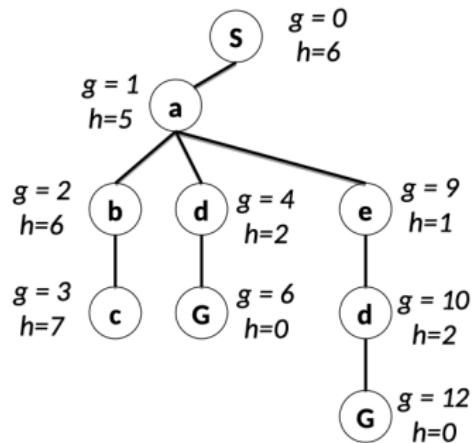
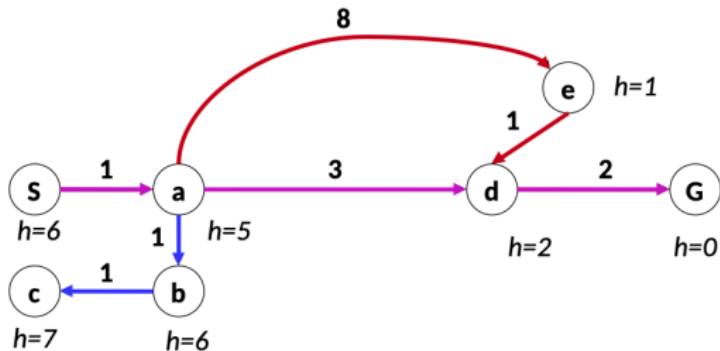


$$\text{Dijkstra (UCS)} \ g(n) + \text{Greedy } h(n) = \text{A}^* \ g(n) + h(n)$$

Example: Teg Grenager

# A\* search

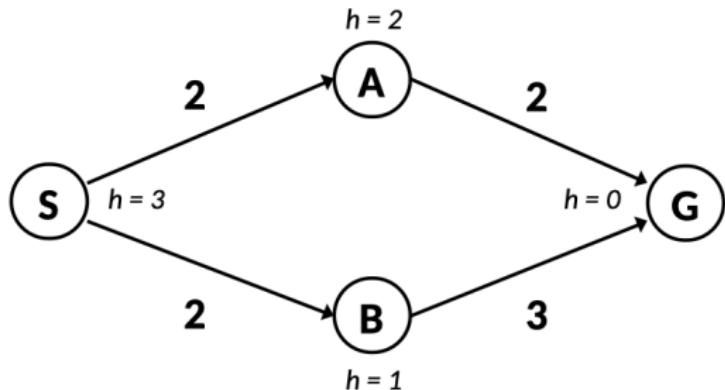
Example



$$\text{Dijkstra (UCS)} \ g(n) + \text{Greedy } h(n) = \text{A}^* \ g(n) + h(n)$$

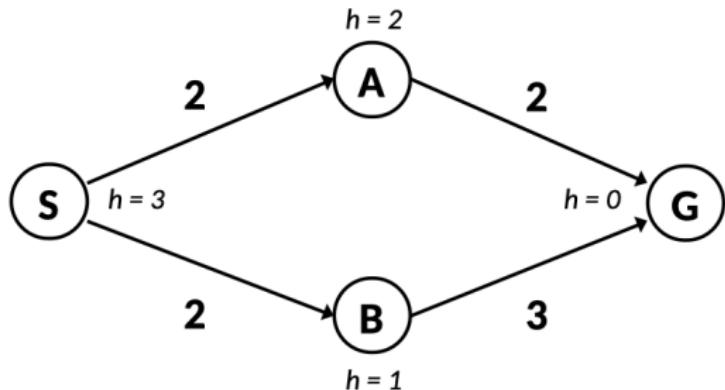
Example: Teg Grenager

## A\* termination



<https://inst.eecs.berkeley.edu/~cs188/su20/>

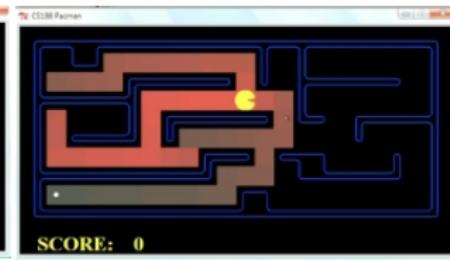
# A\* termination



Stop when dequeue a goal

<https://inst.eecs.berkeley.edu/~cs188/su20/>

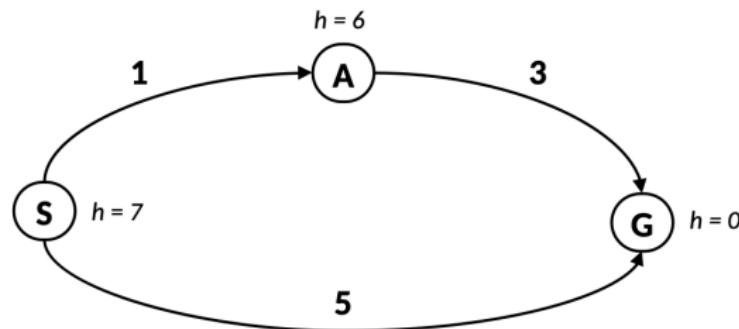
# Greedy vs. UCS vs. A\*



<https://inst.eecs.berkeley.edu/~cs188/su20/>



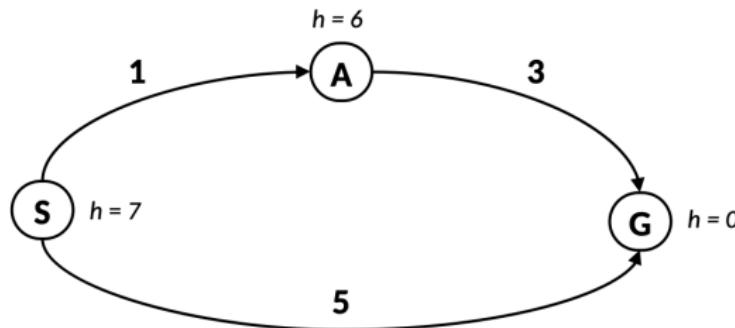
# Interactive block I: Is A\* optimal?



<https://inst.eecs.berkeley.edu/~cs188/su20/>



## Interactive block I: Is A\* optimal?

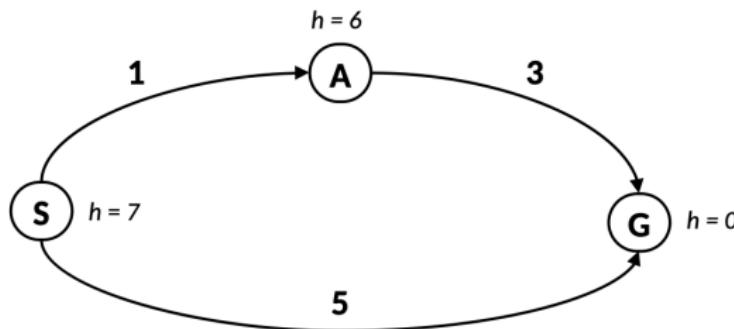


- Actual bad goal cost  $\leq$  estimated good goal cost

<https://inst.eecs.berkeley.edu/~cs188/su20/>



## Interactive block I: Is A\* optimal?



- Actual bad goal cost  $\leq$  estimated good goal cost
- We need estimates to be less than actual costs!

<https://inst.eecs.berkeley.edu/~cs188/su20/>

# Properties of A\* heuristic

Necessary and sufficient for optimality



# Properties of A\* heuristic

Necessary and sufficient for optimality



## Admissible

Never over-estimating the true cost.

# Properties of A\* heuristic

Necessary and sufficient for optimality



**Admissible**

Never over-estimating the true cost.

Sufficient but not necessary condition:

**Consistent**

Never over-estimating the growth of the path cost.

# Properties of A\* heuristic

Necessary and sufficient for optimality



## Admissible

$$0 \leq h(n) \leq h^*(n)$$

# Properties of A\* heuristic

Necessary and sufficient for optimality



Admissible

$$0 \leq h(n) \leq h^*(n)$$

Consistent

$$h(n) \leq h(n') + c(n, n')$$

# Properties of A\* heuristic

Necessary and sufficient for optimality



Admissible

$$0 \leq h(n) \leq h^*(n)$$

Consistent

$$h(n) \leq h(n') + c(n, n')$$

Exercise: prove that consistent implies admissible

# Properties of A\* heuristic

Necessary and sufficient for optimality



## Admissible

$$0 \leq h(n) \leq h^*(n)$$

## Consistent

$$h(n) \leq h(n') + c(n, n')$$

Exercise: prove that consistent implies admissible

Exercise: give a heuristic that is admissible but not consistent



15 minute(s) break!



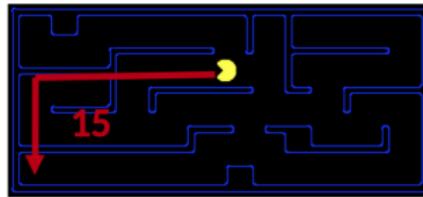
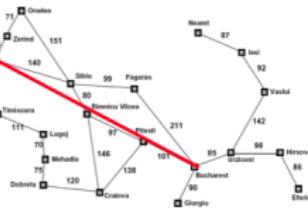
# How to find an admissible heuristic?

This is the challenge



- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics

366



<https://inst.eecs.berkeley.edu/~cs188/su20/>

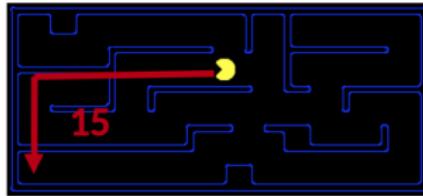
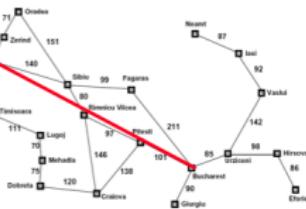
# How to find an admissible heuristic?

This is the challenge



- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to **relaxed** problems, where new actions are available

366



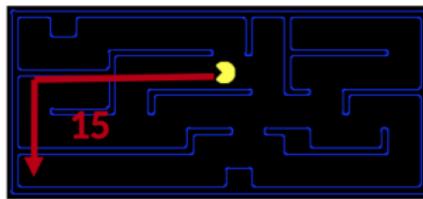
<https://inst.eecs.berkeley.edu/~cs188/su20/>

# How to find an admissible heuristic?

This is the challenge



- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to **relaxed** problems, where new actions are available
- Inadmissible heuristics are often useful too



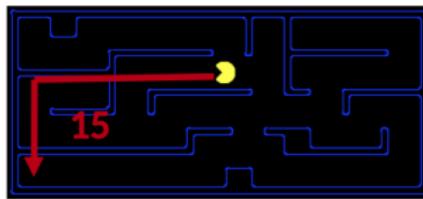
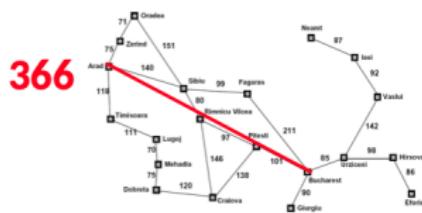
<https://inst.eecs.berkeley.edu/~cs188/su20/>

# How to find an admissible heuristic?

This is the challenge



- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to **relaxed** problems, where new actions are available
- Inadmissible heuristics are often useful too



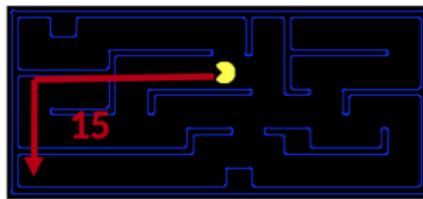
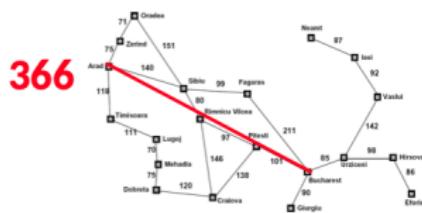
<https://inst.eecs.berkeley.edu/~cs188/su20/>

# How to find an admissible heuristic?

This is the challenge



- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to relaxed problems, where new actions are available
- Inadmissible heuristics are often useful too



<https://inst.eecs.berkeley.edu/~cs188/su20/>



- Tree search:
  - A\* is optimal if heuristic is admissible
  - UCS is a special case ( $h = 0$ )



# Optimality revisited

- Tree search:
  - A\* is optimal if heuristic is admissible
  - UCS is a special case ( $h = 0$ )
- Graph search:
  - A\* optimal if heuristic is consistent
  - UCS optimal ( $h = 0$  is consistent)



# Optimality revisited

- Tree search:
  - A\* is optimal if heuristic is admissible
  - UCS is a special case ( $h = 0$ )
- Graph search:
  - A\* optimal if heuristic is consistent
  - UCS optimal ( $h = 0$  is consistent)
- Consistency implies admissibility, converse is false
- In general, most natural admissible heuristics tend to be consistent, especially if from relaxed problems

# Optimal efficiency



for more details see Dechter & Pearl (1985)



## Optimality

Finds a shortest path.

A\* is optimal if the heuristic is admissible.

for more details see Dechter & Pearl (1985)



# Optimal efficiency

## Optimality

Finds a shortest path.

A\* is optimal if the heuristic is admissible.

## Optimally efficient

Does the least work of any such algorithm.

A\* is optimally efficient if the heuristic is consistent.

for more details see Dechter & Pearl (1985)



## Optimality

Finds a shortest path.

A\* is optimal if the heuristic is admissible.

## Optimally efficient

Does the least work of any such algorithm.

A\* is optimally efficient if the heuristic is consistent.

Efficiency = set of nodes expanded, not number of node expansions

for more details see Dechter & Pearl (1985)



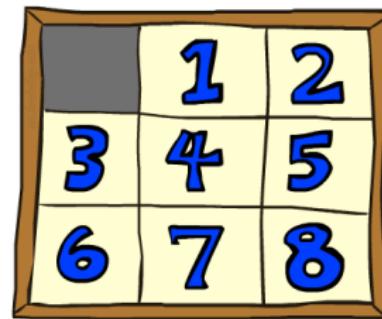
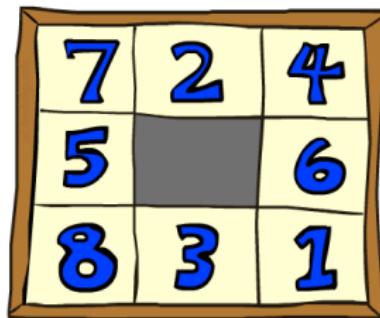
## Interactive block II: Admissible heuristic for sliding puzzle



<https://cs50.harvard.edu/ai/2020>



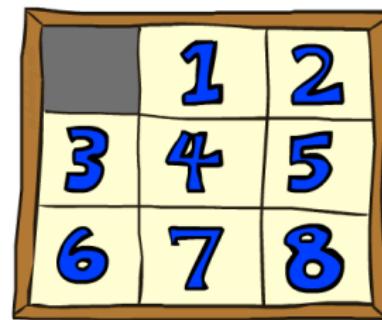
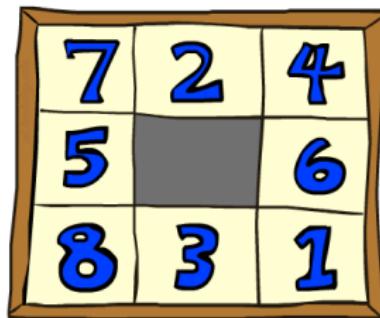
## Interactive block II: Admissible heuristic for sliding puzzle



<https://inst.eecs.berkeley.edu/~cs188/su20/>



## Interactive block II: Admissible heuristic for sliding puzzle

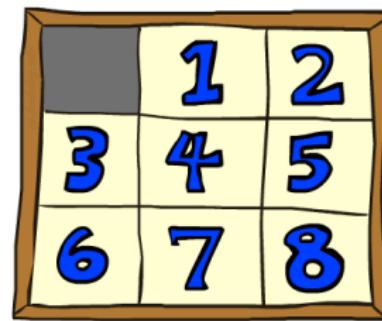
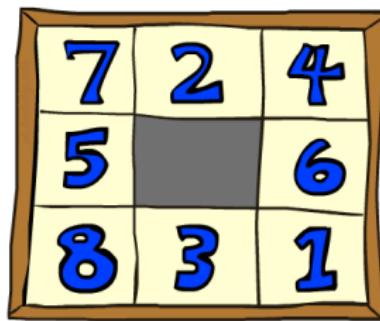


- Number of tiles misplaced

<https://inst.eecs.berkeley.edu/~cs188/su20/>



## Interactive block II: Admissible heuristic for sliding puzzle

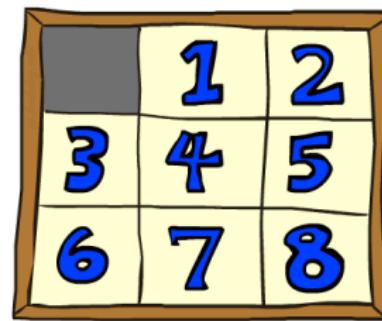
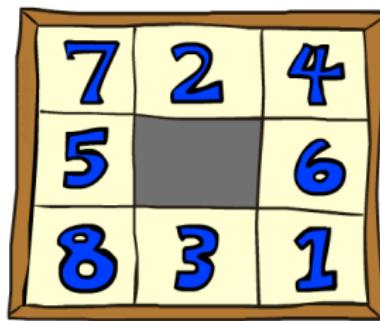


- Number of tiles misplaced
- Total Manhattan distance

<https://inst.eecs.berkeley.edu/~cs188/su20/>



## Interactive block II: Admissible heuristic for sliding puzzle

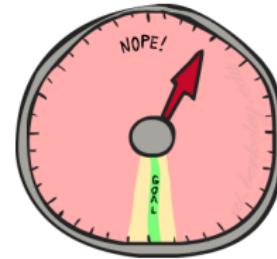
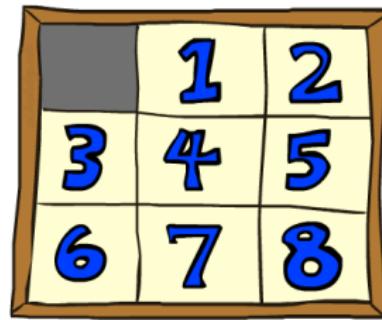
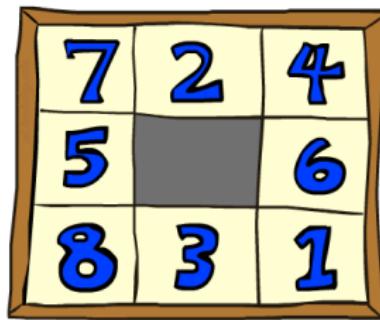


- Number of tiles misplaced
- Total Manhattan distance
- Actual cost to solution?

<https://inst.eecs.berkeley.edu/~cs188/su20/>



## Interactive block II: Admissible heuristic for sliding puzzle



<https://inst.eecs.berkeley.edu/~cs188/su20/>



# A\* extensions

- Anytime A\*
- Block A\*
- D\*
- Field D\*
- Fringe
- Fringe Saving A\* (FSA\*)
- Generalised Adaptive A\* (GAA\*)
- Incremental heuristic search
- Reduced A\*
- **Iterative deepening A\* (IDA\*)**
- Jump point search
- Lifelong Planning A\* (LPA\*)
- New Bidirectional A\* (NBA\*)
- Simplified Memory bounded A\* (SMA\*)
- Theta\*



## Next lecture

- Today we saw A\* search



- Today we saw A\* search
  - Sorry, no OXO today ☺



## Next lecture

- Today we saw A\* search
  - Sorry, no OXO today ☺
  - A\* is optimal + optimally efficient, but space...



## Next lecture

- Today we saw A\* search
  - Sorry, no OXO today ☺
  - A\* is optimal + optimally efficient, but space...
- Next time: adversarial search



## Next lecture

- Today we saw A\* search
  - Sorry, no OXO today ☺
  - A\* is optimal + optimally efficient, but space...
- Next time: adversarial search
- Preparation: see Brightspace



## Next lecture

- Today we saw A\* search
  - Sorry, no OXO today ☺
  - A\* is optimal + optimally efficient, but space...
- Next time: adversarial search
- Preparation: see Brightspace
- Lab on Monday



## Next lecture

- Today we saw A\* search
  - Sorry, no OXO today ☺
  - A\* is optimal + optimally efficient, but space...
- Next time: adversarial search
- Preparation: see Brightspace
- Lab on Monday
- Assignment 1 due: see Brightspace



- Today we saw A\* search
  - Sorry, no OXO today ☺
  - A\* is optimal + optimally efficient, but space...
- Next time: adversarial search
- Preparation: see Brightspace
- Lab on Monday
- Assignment 1 due: see Brightspace
- Questions? TUD Answers



# What is still unclear?

After every lecture...

Give us some homework:  
tell us what is still unclear.

# ALGORITHMS FOR NP-HARD PROBLEMS

## Part I – Heuristic Search

### Lecture 3: Adversarial Search

Neil Yorke-Smith

[n.yorke-smith@tudelft.nl](mailto:n.yorke-smith@tudelft.nl)



©2021–23 — CC BY-NC-NA 4.0

February 2023



# Where are we?

- Last time: A\* search
- Today: minimax search



# Where are we?

- Last time: A\* search
- Today: minimax search
- You prepared:
  - Paper about Deep Blue



# Adversary!

Regular search

Find an answer to a question.

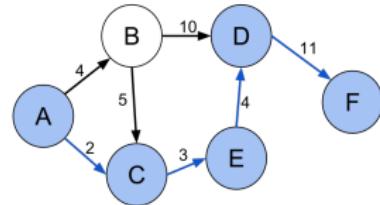


Photo by Mesh on Unsplash



# Adversary!

Regular search

Find an answer to a question.

Adversarial search

But now an opponent that tries to achieve the opposite goal.



Photo by Mesh on Unsplash

# Adversary!

Regular search

Find an answer to a question.

Adversarial search

But now an opponent that tries to achieve the opposite goal.

Two player, sequential deterministic moves, perfect information, discrete zero-sum



Photo by Mesh on Unsplash

# Terminology of games

- single/**two**/multi-player game
  - Solitaire/Chess/Settlers-of-Catan

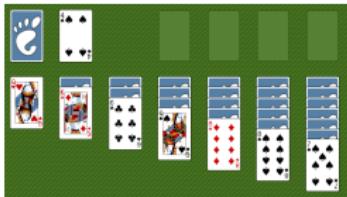


Image sources: [wikimedia.org](https://commons.wikimedia.org)

# Terminology of games

- single/**two**/multi-player game
  - Solitaire/Chess/Settlers-of-Catan
- simultaneous/**sequential** moves
  - Rock-paper-scissors/Go

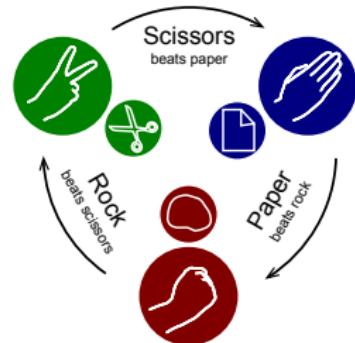


Image sources: [wikimedia.org](#)

# Terminology of games

- single/**two**/multi-player game
  - Solitaire/Chess/Settlers-of-Catan
- simultaneous/**sequential** moves
  - Rock-paper-scissors/Go
- stochastic/**deterministic** moves
  - Poker/Chess



Image sources: [wikimedia.org](#)

# Terminology of games

- single/**two**/multi-player game
  - Solitaire/Chess/Settlers-of-Catan
- simultaneous/**sequential** moves
  - Rock-paper-scissors/Go
- stochastic/**deterministic** moves
  - Poker/Chess
- partial/**perfect** information
  - Card-games/Connect-four



Image sources: [wikimedia.org](#)

# Terminology of games



- single/**two**/multi-player game
  - Solitaire/Chess/Settlers-of-Catan
- simultaneous/**sequential** moves
  - Rock-paper-scissors/Go
- stochastic/**deterministic** moves
  - Poker/Chess
- partial/**perfect** information
  - Card-games/Connect-four
- **discrete**/continual state/actions/time
  - turn-based/real-time strategy



Image sources: wikipedia.org, Samvelyan et al. (2019)

# Terminology of games



- single/**two**/multi-player game
  - Solitaire/Chess/Settlers-of-Catan
- simultaneous/**sequential** moves
  - Rock-paper-scissors/Go
- stochastic/**deterministic** moves
  - Poker/Chess
- partial/**perfect** information
  - Card-games/Connect-four
- **discrete**/continual state/actions/time
  - turn-based/real-time strategy
- **zero-sum**/general-sum game
  - Chess/Settlers-of-Catan



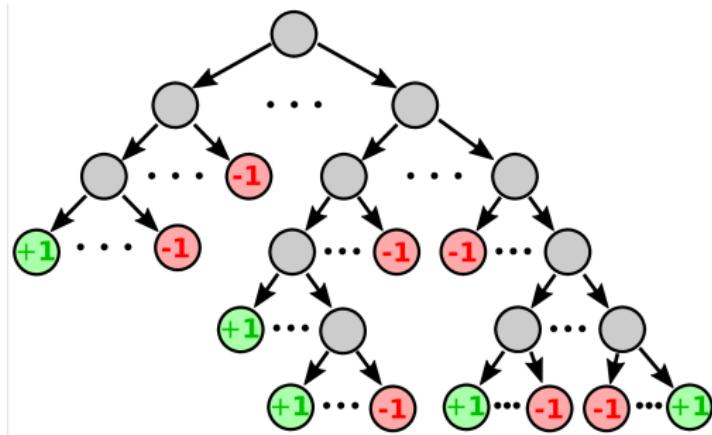
Image sources: [wikimedia.org](#)

# Minimax evaluation

First, search trees as we did so far



Suppose value  $+1$  for a win,  $-1$  for a lose

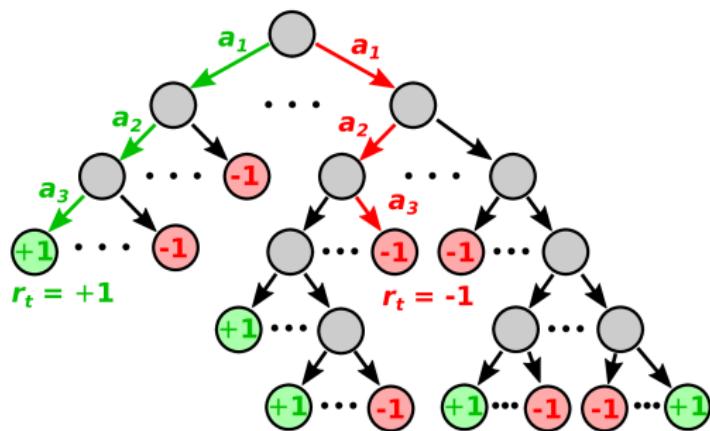


# Minimax evaluation

First, search trees as we did so far

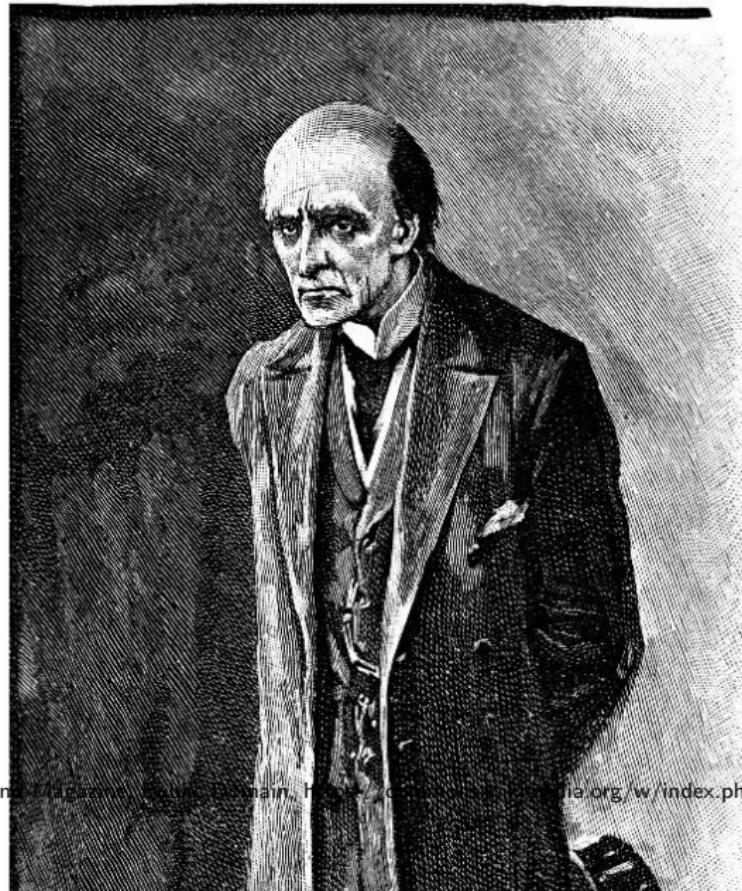


Winning path, losing path



# Minimax evaluation

## Adversary moves



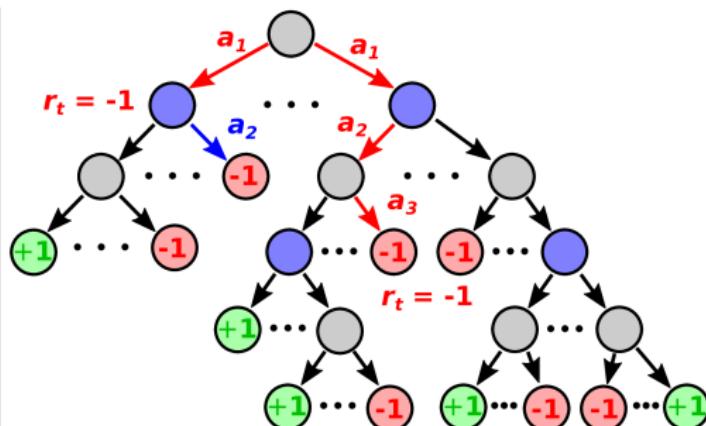
By Sidney Paget - The Strand Magazine, Public Domain, <https://commons.wikimedia.org/w/index.php?curid=69209162>

# Minimax evaluation

Adversary moves



Adversary's plays  $\implies$  left-hand  $a_1$  is now a losing move!

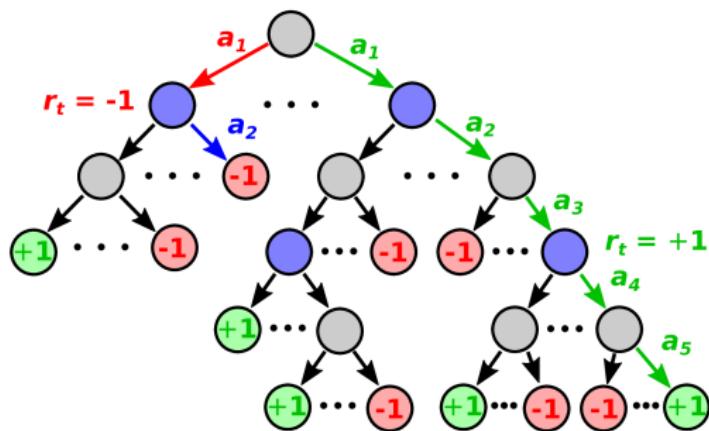


# Minimax evaluation

Adversary moves



So our moves must account for his. How?

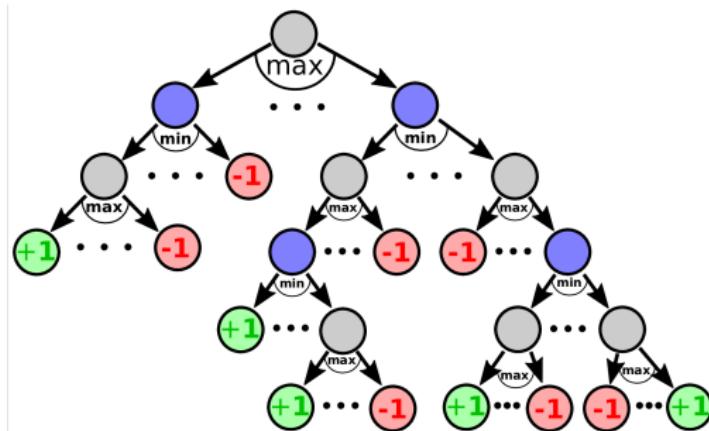


# Minimax evaluation

Max for me, min for him



**Key idea:** play value-maximising moves, assuming that adversary plays value-minimising moves.

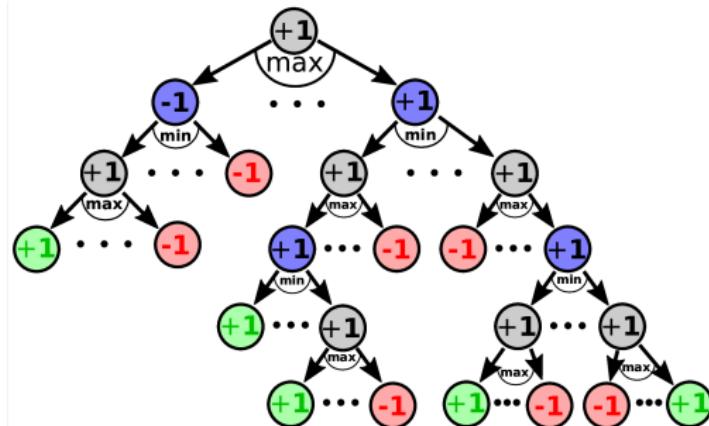


# Minimax evaluation

Max for me, min for him



At each node, compute minimax score from node's children

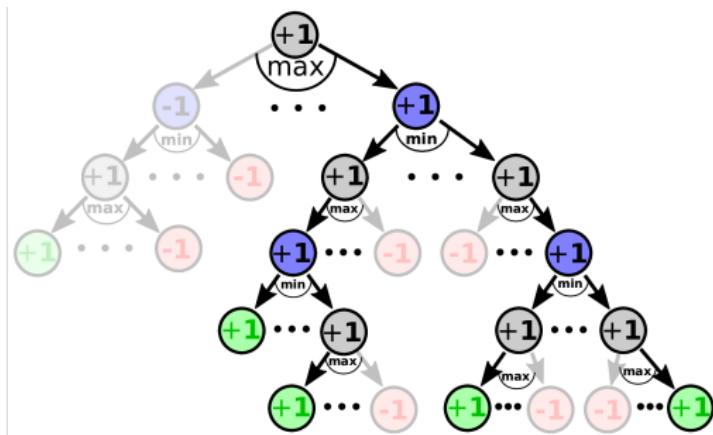


# Minimax evaluation

Max for me, min for him

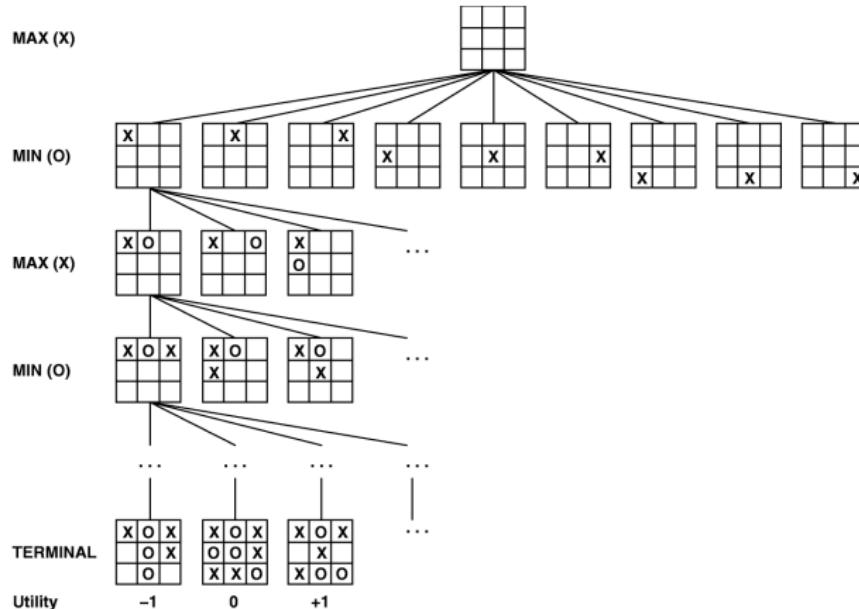


Choose moves that lead to maximum score



# Search tree for OXO

Remember this?



<https://www.massey.ac.nz/~mjjohnso/>



Suppose  $X$  goes first, in the top-left corner

Write down the whole search tree, assuming an adversary who plays optimally.

# Improving minimax

What now?



Minimax search doesn't scale

Works for OXO, not for chess.

# Improving minimax

What now?



Minimax search doesn't scale

Works for OXO, not for chess. Or shogi, or go, or poker, ...

# Improving minimax

What now?



Minimax search doesn't scale

Works for OXO, not for chess. Or shogi, or go, or poker, ...

## Ideas

- Don't fathom the tree – depth-limited

# Improving minimax

What now?



Minimax search doesn't scale

Works for OXO, not for chess. Or shogi, or go, or poker, ...

## Ideas

- Don't fathom the tree – depth-limited
- Prune useless subtrees – alpha-beta pruning

# Improving minimax

What now?



Minimax search doesn't scale

Works for OXO, not for chess. Or shogi, or go, or poker, ...

## Ideas

- Don't fathom the tree – depth-limited
- Prune useless subtrees – alpha-beta pruning
- Search stochastically – MCTS Lecture 4

# Depth-limited minimax search

Don't explore the whole tree



```
fun minimax(n: node): int =  
    if leaf(n) then return evaluate(n)  
    if n is a max node  
        v := L  
        for each child of n  
            v' := minimax (child)  
            if v' > v, v:= v'  
        return v  
    if n is a min node  
        v := W  
        for each child of n  
            v' := minimax (child)  
            if v' < v, v:= v'  
        return v
```

# Depth-limited minimax search

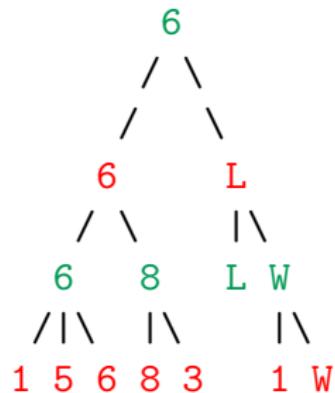
Don't explore the whole tree



```
fun minimax(n: node, d: int): int =  
    if leaf(n) or d=0 return evaluate(n)  
    if n is a max node  
        v := L  
        for each child of n  
            v' := minimax (child,d-1)  
            if v' > v, v:= v'  
        return v  
    if n is a min node  
        v := W  
        for each child of n  
            v' := minimax (child,d-1)  
            if v' < v, v:= v'  
        return v
```

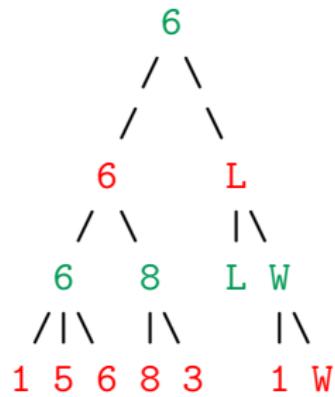
# Depth-limited minimax search

Don't explore the whole tree



# Depth-limited minimax search

Don't explore the whole tree



How deep to search?



15 minute(s) break!



# Improving minimax

What now?



Minimax search doesn't scale

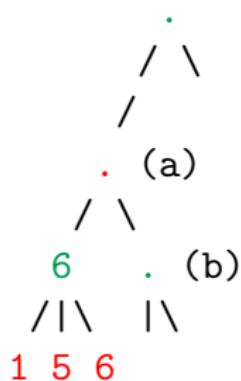
Works for OXO, not for chess. Or shogi, or go, or poker, ...

## Ideas

- Don't fathom the tree – depth-limited
- Prune useless subtrees – alpha-beta pruning
- Search stochastically – MCTS Lecture 4

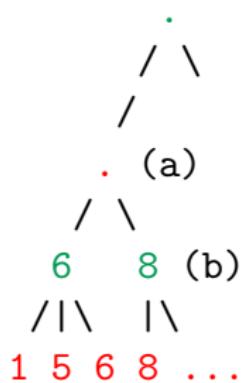
# Alpha-beta pruning

Prune useless subtrees



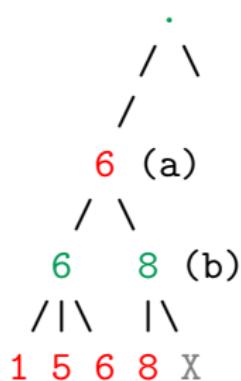
# Alpha-beta pruning

Prune useless subtrees



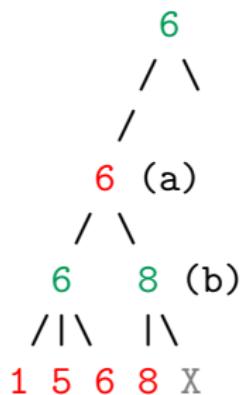
# Alpha-beta pruning

Prune useless subtrees



# Alpha-beta pruning

Prune useless subtrees



# Alpha-beta pruning

Prune useless subtrees



```
fun minimax(n: node, d: int, min: int, max: int): int =  
    if leaf(n) or d=0 return evaluate(n)  
    if n is a max node  
        v := min  
        for each child of n  
            v' := minimax (child,d-1,...,...)  
            if v' > v, v:= v'  
            if v > max return max  
        return v  
    if n is a min node  
        v := max  
        for each child of n  
            v' := minimax (child,d-1,...,...)  
            if v' < v, v:= v'  
            if v < min return min  
    return v
```

# Alpha-beta pruning

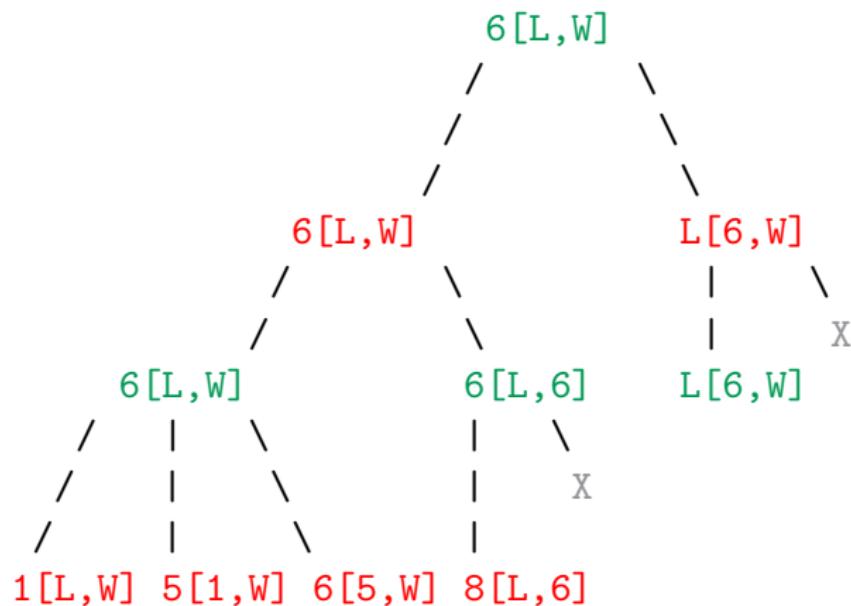
Prune useless subtrees



```
fun minimax(n: node, d: int, min: int, max: int): int =  
    if leaf(n) or d=0 return evaluate(n)  
    if n is a max node  
        v := min  
        for each child of n  
            v' := minimax (child,d-1,v,max)  
            if v' > v, v:= v'  
            if v > max return max  
        return v  
    if n is a min node  
        v := max  
        for each child of n  
            v' := minimax (child,d-1,min,v)  
            if v' < v, v:= v'  
            if v < min return min  
    return v
```

# Alpha-beta pruning

Prune useless subtrees



# Alpha-beta pruning

Prune useless subtrees



How to prune effectively?

# Alpha-beta pruning

Prune useless subtrees



## How to prune effectively?

- Order well the nodes (high max, low min)

# Alpha-beta pruning

Prune useless subtrees



## How to prune effectively?

- Order well the nodes (high max, low min)
- Previous (partial) searches of the tree

# Alpha-beta pruning

Prune useless subtrees

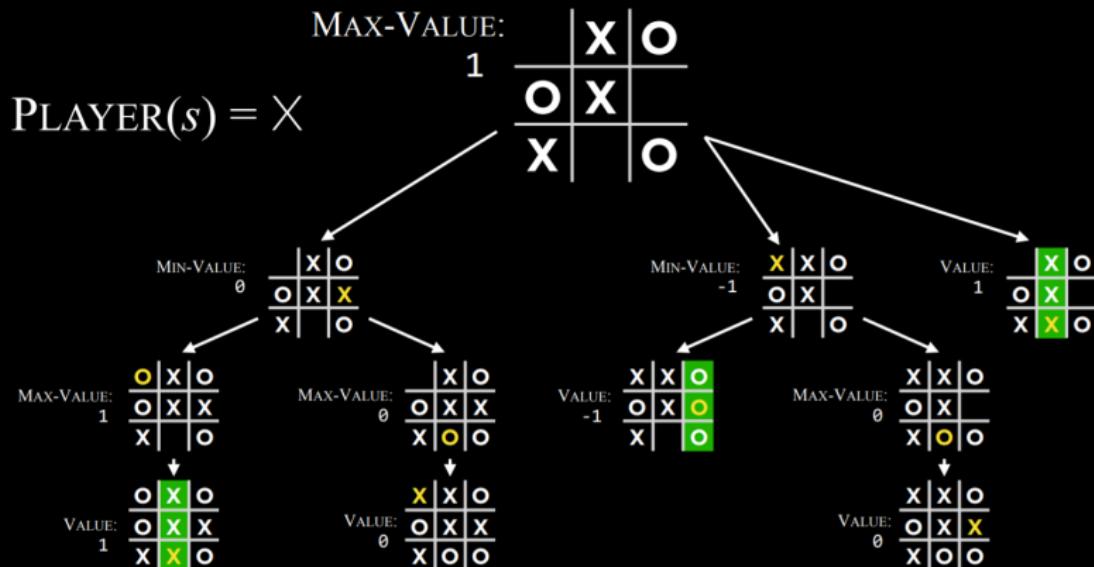


## How to prune effectively?

- Order well the nodes (high max, low min)
- Previous (partial) searches of the tree
- Good heuristic evaluation function



## Interactive block II: Does $\alpha$ - $\beta$ help here?



<https://cs50.harvard.edu/ai/2020>

# Example: Deep Blue

Famous in 1997



By James the photographer - <https://www.flickr.com/photos/james3761@N00/592436598/>, CC BY 2.0,  
<https://commons.wikimedia.org/w/index.php?curid=3511068>



Evaluation functions are subtle

Hard to code a good evaluation for complex states.



Evaluation functions are subtle

Hard to code a good evaluation for complex states.

Example: Chess

Material + movement + king safety + passed pawns + ...



Evaluation functions are subtle

Hard to code a good evaluation for complex states.

Example: Chess

Material + movement + king safety + passed pawns + ...

Idea: Machine learning

Efficient neural networks (NNUE).

# Example: Stockfish 12

Famous in 2020



Stockfish 12 has added neural network evaluation abilities.

As of October 2020, Stockfish is the highest-rated engine according to the computer chess rating list (CCRL) with a rating of 3514—it is the only engine with a rating above 3500. According to the July 2020 Swedish Chess Computer Association (SSDF) rating list, Stockfish 9 is ranked #3, Stockfish 10 is ranked #2, and Stockfish 11 is ranked #1 with a rating of 3558. Taking the top three spots with three different versions is quite impressive.

According to this great video on the strongest chess engines of all time (based on the SSDF rating lists), Stockfish is the strongest engine of all time—a sentiment that is widely shared in the chess community.

<https://www.chess.com/terms/stockfish-chess-engine>



- Today we saw minimax search



- Today we saw minimax search
  - With enhancements and NN evaluation = state of the art for chess



- Today we saw minimax search
  - With enhancements and NN evaluation = state of the art for chess
- Next time: monte carlo tree search



- Today we saw minimax search
  - With enhancements and NN evaluation = state of the art for chess
- Next time: monte carlo tree search
  - With enhancements and NN evaluation = state of the art for go



- Today we saw minimax search
  - With enhancements and NN evaluation = state of the art for chess
- Next time: monte carlo tree search
  - With enhancements and NN evaluation = state of the art for go
- Preparation: see Brightspace



- Today we saw minimax search
  - With enhancements and NN evaluation = state of the art for chess
- Next time: monte carlo tree search
  - With enhancements and NN evaluation = state of the art for go
- Preparation: see Brightspace
- Lab on Monday



- Today we saw minimax search
  - With enhancements and NN evaluation = state of the art for chess
- Next time: monte carlo tree search
  - With enhancements and NN evaluation = state of the art for go
- Preparation: see Brightspace
- Lab on Monday
- Assignment 1 due: see Brightspace



- Today we saw minimax search
  - With enhancements and NN evaluation = state of the art for chess
- Next time: monte carlo tree search
  - With enhancements and NN evaluation = state of the art for go
- Preparation: see Brightspace
- Lab on Monday
- Assignment 1 due: see Brightspace
- Questions? TUD Answers



# What is still unclear?

After every lecture...

Give us some homework:  
tell us what is still unclear.



# References I

Mikayel Samvelyan, Tabish Rashid, Christian Schroeder de Witt, Gregory Farquhar, Nantas Nardelli, Tim G. J. Rudner, Chia-Man Hung, Philip H. S. Torr, Jakob Foerster, and Shimon Whiteson. The StarCraft Multi-Agent Challenge. *CoRR*, abs/1902.04043, 2019.

# ALGORITHMS FOR NP-HARD PROBLEMS

## Part I – Heuristic Search

### Lecture 4: Monte Carlo Tree Search

Wendelin Böhmer and Neil Yorke-Smith

n.yorke-smith@tudelft.nl



©2021–23 — CC BY-NC-NA 4.0

February 2023



# Where are we?

- Last time: minimax search
- Today: exploiting stochasticity



# Where are we?

- Last time: minimax search
- Today: exploiting stochasticity
- You prepared:
  - Paper about AlphaGo

# Terminology of games



- single/**two**/multi-player game
  - Solitaire/Chess/Settlers-of-Catan
- simultaneous/**sequential** moves
  - Rock-paper-scissors/Go
- stochastic/**deterministic** moves
  - Poker/Chess
- partial/**perfect** information
  - Card-games/Connect-four
- **discrete**/continual state/actions/time
  - turn-based/real-time strategy
- **zero-sum**/general-sum game
  - Chess/Settlers-of-Catan



Image sources: [wikimedia.org](#)

# Formal definition

## Markov Decision Process



- Deterministic episodic two-player MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{T}, \mathcal{R} \rangle$ 
  - $s \in \mathcal{S}$  denotes all game-states [discrete, perfect information]
  - $a \in \mathcal{A}(s)$  are all actions that can be played in  $s$  [discrete, sequential]
  - $\mathcal{P}(s, a) = s'$  determines next state  $s' \in \mathcal{S}$  [deterministic moves]
  - $\mathcal{T}(s) \in \{\top, \perp\}$  whether the game terminates (ends) in  $s$  [episodic]
  - $\mathcal{R}(s) \in \mathbb{R}$  returns the reward for terminal state  $s$
  - $p(s) \in \{-1, 1\}$  returns which player's move it is [two sequential players]
- Note that we only consider *reward* at the end of the game
  - +1 for winning, 0 for draw, -1 for loss
  - other definitions are possible, but can change solution

# State interface (1)



- two-player (1 vs. -1) zero-sum games with perfect information

```
1 class State:  
2     def __init__(self, player=1):  
3         self.player = player          #  $p(s) \in \{-1, 1\}$   
4  
5     def actions(self):             #  $\mathcal{A}(s)$   
6         return []      # list of legal actions  
7  
8     def transition(self, action):   #  $\mathcal{P}(s,a) \in \mathcal{S}$   
9         return None    # new state after action  
10  
11    def terminal(self):           #  $\mathcal{T}(s) \in \{\top, \perp\}$   
12        return None    # boolean whether the game ends  
13  
14    def reward(self):            #  $\mathcal{R}(s) \in \mathbb{R}$   
15        return None    # 1 for won, 0 for draw and -1 for lost
```



# Action interface

- two-player (1 vs. -1) zero-sum games with perfect information
- discrete sequential deterministic moves

```
1 class Action:  
2     def __init__(self, player, x, ...):  
3         self.player = player  
4         self.x = x  
5         # more parameters describing the action  
6  
7     def __str__(self):          # to look better  
8         return str(self.x)  
9  
10    def __eq__(self, other): # for comparisons ('if a in list')  
11        return self.player == other.player and self.x == other.x  
12  
13    def __hash__(self):      # for use in dictionaries (hash maps)  
14        return hash((self.player, self.x))
```

# Example: Noughts-and-crosses (OXO)



```
1 class NoughtsAndCrosses (State):
2     def __init__(self, player=1):
3         super().__init__(player=player)
4         self.board = [[0 for _ in range(3)] for _ in range(3)]
5
6     def actions(self):           #  $\mathcal{A}(s)$ 
7         # returns a list with an action for each 0 in self.board
8
9     def transition(self, action):    #  $\mathcal{P}(s,a) \in \mathcal{S}$ 
10    state = deepcopy(self)
11    state.board[action.x][action.y] = action.player
12    state.player *= -1
13    return state
14
15    def reward(self):            #  $\mathcal{R}(s) \in \mathbb{R}$ 
16        # checks all rows, columns and diagonals <x>:
17        # return +1 if sum(x) == +3
18        # return -1 if sum(x) == -3
19        # return 0 if no x triggered a return
20
21    def terminal(self):          #  $\mathcal{T}(s) \in \{\top, \perp\}$ 
22    return self.reward() != 0 or len(self.actions()) == 0
```



# Interactive block I: two-player games



Multiple (or no) answers may be correct!

- How many distinct states has noughts-and-crosses (OXO)?
  - (A)  $n = 2^9$
  - (B)  $n = 3^9$
  - (C)  $n = 9!$
- Which *properties* has the game backgammon?
  - (E) simultaneous moves
  - (F) deterministic moves
  - (G) perfect information
  - (H) continuous time
  - (I) discrete states



image source wikipedia.org



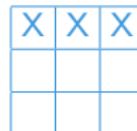
# Interactive block I: two-player games



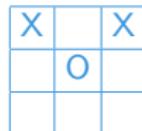
Multiple (or no) answers may be correct!

- How many distinct states has noughts-and-crosses (OXO)?

- A  $n = 2^9$
- B  $n = 3^9$
- C  $n = 9!$



not reachable  
 $\Rightarrow n \neq 2^9$



reachable by two paths  
 $\Rightarrow n \neq 9!$

- Which *properties* has the game backgammon?

- E simultaneous moves
- F deterministic moves
- G perfect information ✓
- H continuous time
- I discrete states ✓



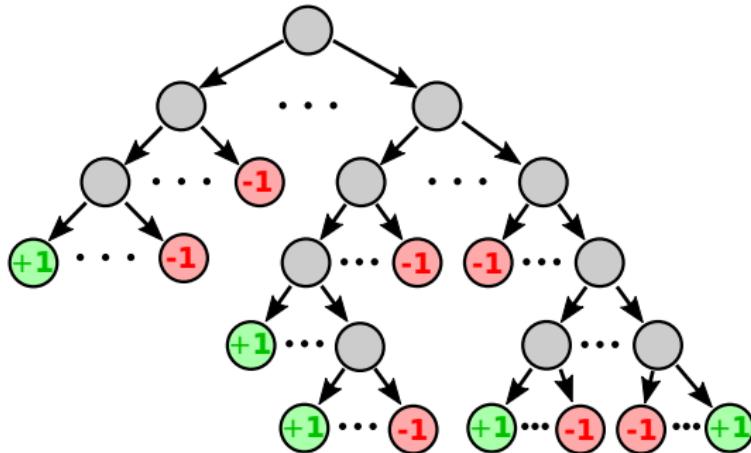
image source [wikimedia.org](#)

# Decision trees (1)

Recall last lecture



- Find a sequence of actions  $a_1, \dots, a_n$  that leads to a win
  - standard search problem for single player games
  - many possible solutions: easy problem

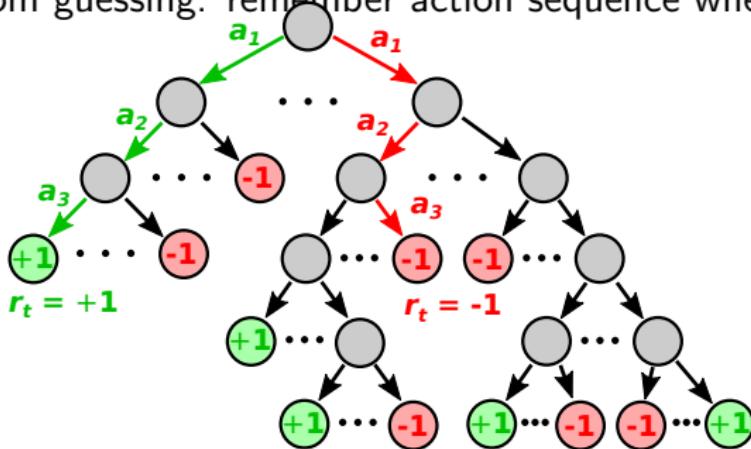


# Decision trees (1)

Recall last lecture



- Find a sequence of actions  $a_1, \dots, a_n$  that leads to a win
  - standard search problem for single player games
  - many possible solutions: easy problem
- Random guessing: remember action sequence when  $r = +1$



# Decision trees (1)

Recall last lecture



- Find a sequence of actions  $a_1, \dots, a_n$  that leads to a win
  - standard search problem for single player games
  - many possible solutions: easy problem
- Random guessing: remember action sequence when  $r = +1$

```
1 def random_guessing(initial_state):
2     while True:
3         state, actions = initial_state, []
4         while not state.terminal():
5             actions.append(random.choice(state.actions()))
6             state = state.transition(actions[-1])
7         if state.reward() == +1:
8             return actions
```

# Decision trees (1)

Recall last lecture



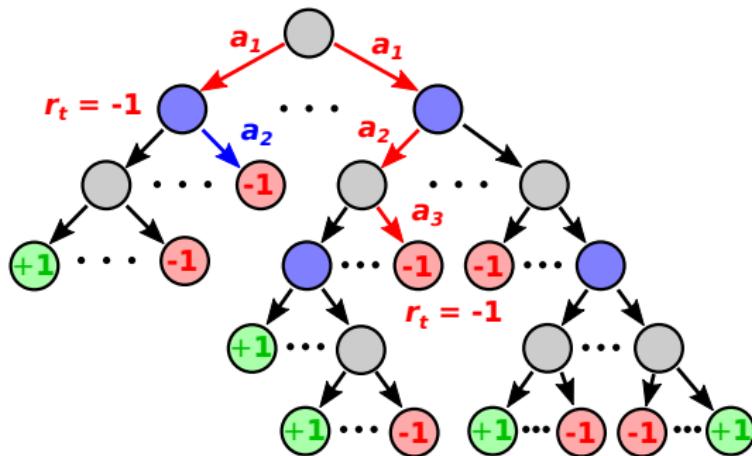
- Find a sequence of actions  $a_1, \dots, a_n$  that leads to a win
  - standard search problem for single player games
  - many possible solutions: easy problem
- Random guessing: remember action sequence when  $r = +1$ 
  - random execution time

# Decision trees (2)

Recall last lecture



- But the **opponent** chooses some of the actions!
  - opponent can easily prevent random solutions

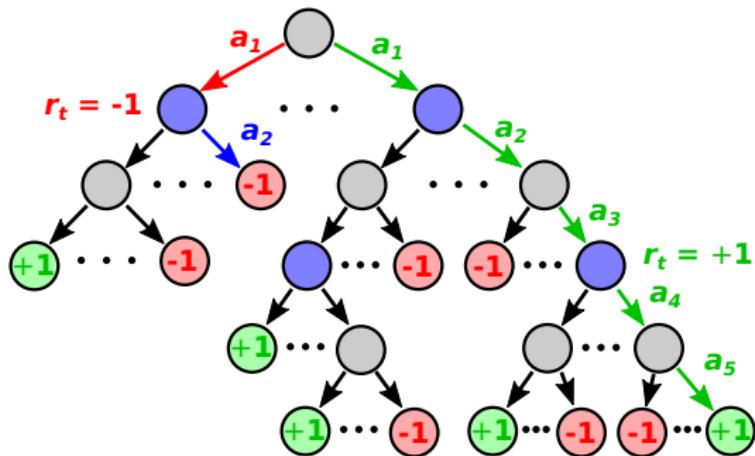


# Decision trees (2)

Recall last lecture



- But the **opponent** chooses some of the actions!
  - opponent can easily prevent random solutions
  - in fact, all winning action sequence in this tree ...

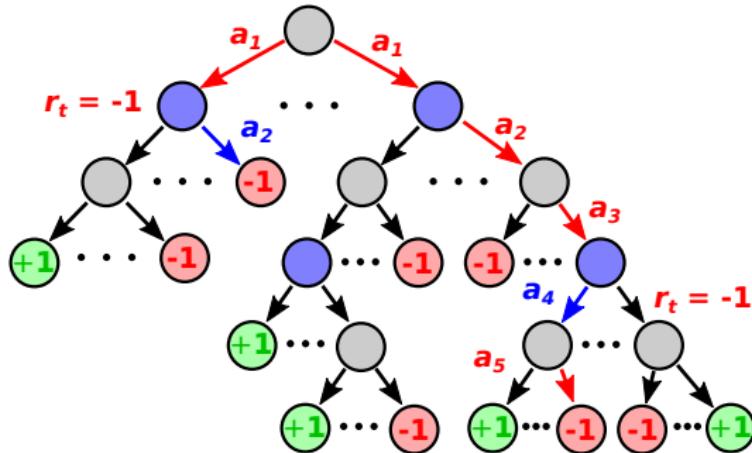


# Decision trees (2)

Recall last lecture



- But the **opponent** chooses some of the actions!
  - opponent can easily prevent random solutions
  - in fact, all winning action sequence in this tree ...
  - ... that can be easily disturbed





- Problem very complex for *general-sum* games
  - sometimes cooperation, sometimes competition
  - very complex Nash equilibria ( Game Theory)



- Zero-sum: “**one players' gain is the other's loss**”
  - every reward  $r$  for one player is a punishment  $-r$  for the other
  - two-player zero sum game can be expressed by *one* function
  - player 1 *maximizes* function, player -1 *minimizes* function

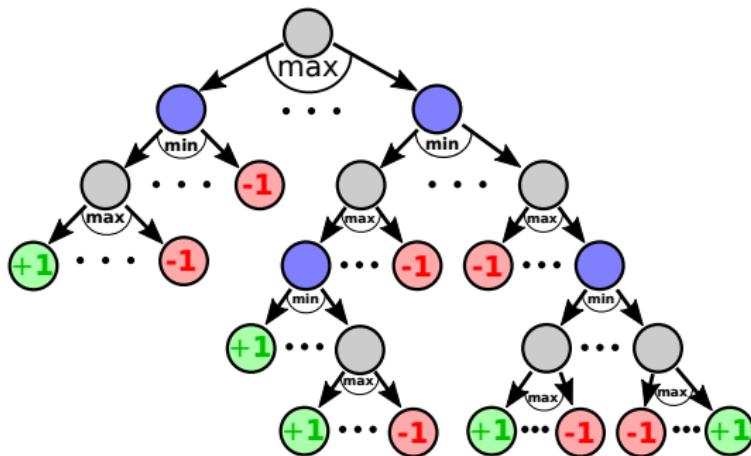
Image sources: wikimedia.org

# Minmax value function

Formalised in MDP framework



- The value function  $V(s)$  of a state  $s \in \mathcal{S}$ 
  - maximizes return for  $p(s) = 1$  (grey nodes)
  - minimizes return for  $p(s) = -1$  (blue nodes)

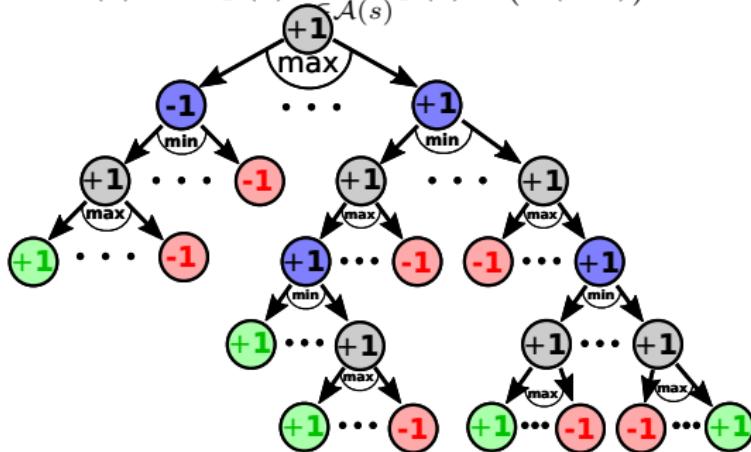


# Minmax value function

Formalised in MDP framework



- The value function  $V(s)$  of a state  $s \in \mathcal{S}$ 
  - maximizes return for  $p(s) = 1$  (grey nodes)
  - minimizes return for  $p(s) = -1$  (blue nodes)
  - if  $\mathcal{T}(s)$  then  $V(s) := \mathcal{R}(s)$
  - else  $V(s) := p(s) \max_{\mathcal{A}(s)} p(s) V(\mathcal{P}(s, a))$

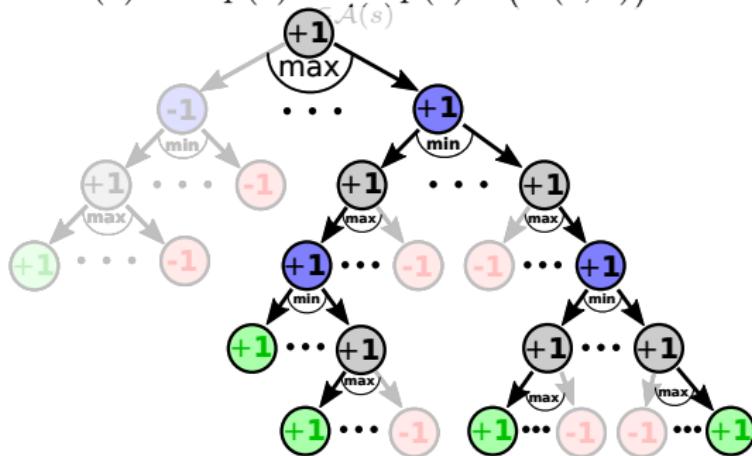


# Minmax value function

Formalised in MDP framework



- The value function  $V(s)$  of a state  $s \in \mathcal{S}$ 
  - maximizes return for  $p(s) = 1$  (grey nodes)
  - minimizes return for  $p(s) = -1$  (blue nodes)
  - if  $\mathcal{T}(s)$  then  $V(s) := \mathcal{R}(s)$
  - else  $V(s) := p(s) \max_{A(s)} p(s) V(\mathcal{P}(s, a))$



# Decision tree implementation



```
1 class Node:
2     def __init__(self, state):
3         self.state = state                      #  $s \in \mathcal{S}$ 
4         self.children = {}                      #  $\{a : \mathcal{P}(s, a), \dots\}$ 
5         self.total_value = 0                   #  $\tilde{V}(s) = \frac{1}{N(s)} \sum_{t=1}^{N(s)} r_t$ 
6         self.num_visits = 0                   #  $N(s)$ 
7
8 class DecisionTree:
9     def __init__(self, initial_state):
10        self.root = Node(initial_state) # root of the tree
11
12    def node_value(self, node, parent=None, **kwargs):
13        value = node.total_value / node.num_visits
14        if parent is not None:
15            value *= parent.state.player
16        return value                         #  $p(\text{parent}.s) \tilde{V}(\text{node}.s)$ 
```



## Interactive block II: zero-sum games



Multiple (or no) answers may be correct!

Let the **Minmax-value** for this question be defined like this:

$$V(s) = \begin{cases} \alpha \mathcal{R}(s) - \beta & , \text{if } \mathcal{T}(s) \\ \gamma p(s) \max_{a \in \mathcal{A}(s)} p(s) V(\mathcal{P}(s, a)) - \delta & , \text{if } \neg \mathcal{T}(s) \end{cases}$$

- Which combination(s) will return the **shortest winning path**?

A)  $\alpha = 1, \beta = 0, \gamma = 1, \delta = 0$

B)  $\alpha = \frac{1}{2}, \beta = 0, \gamma = 1, \delta = 0$

C)  $\alpha = 1, \beta = \frac{1}{2}, \gamma = 1, \delta = 0$

D)  $\alpha = \frac{1}{2}, \beta = \frac{1}{2}, \gamma = 1, \delta = 0$

E)  $\alpha = 1, \beta = 0, \gamma = \frac{1}{2}, \delta = 0$

F)  $\alpha = \frac{1}{2}, \beta = 0, \gamma = \frac{1}{2}, \delta = 0$

G)  $\alpha = 1, \beta = \frac{1}{2}, \gamma = \frac{1}{2}, \delta = 0$

H)  $\alpha = \frac{1}{2}, \beta = \frac{1}{2}, \gamma = \frac{1}{2}, \delta = 0$

I)  $\alpha = 1, \beta = 0, \gamma = 1, \delta = \frac{1}{2}$

J)  $\alpha = \frac{1}{2}, \beta = 0, \gamma = 1, \delta = \frac{1}{2}$

K)  $\alpha = 1, \beta = \frac{1}{2}, \gamma = 1, \delta = \frac{1}{2}$

L)  $\alpha = \frac{1}{2}, \beta = \frac{1}{2}, \gamma = 1, \delta = \frac{1}{2}$

M)  $\alpha = 1, \beta = 0, \gamma = \frac{1}{2}, \delta = \frac{1}{2}$

N)  $\alpha = \frac{1}{2}, \beta = 0, \gamma = \frac{1}{2}, \delta = \frac{1}{2}$

O)  $\alpha = 1, \beta = \frac{1}{2}, \gamma = \frac{1}{2}, \delta = \frac{1}{2}$

P)  $\alpha = \frac{1}{2}, \beta = \frac{1}{2}, \gamma = \frac{1}{2}, \delta = \frac{1}{2}$



Multiple (or no) answers may be correct!

Let the **Minmax-value** for this question be defined like this:

$$V(s) = \begin{cases} \alpha \mathcal{R}(s) - \beta & , \text{if } \mathcal{T}(s) \\ \gamma p(s) \max_{a \in \mathcal{A}(s)} p(s) V(\mathcal{P}(s, a)) - \delta & , \text{if } \neg \mathcal{T}(s) \end{cases}$$

- Which combination(s) will return the **shortest winning path**?

A)  $\alpha = 1, \beta = 0, \gamma = 1, \delta = 0$

B)  $\alpha = \frac{1}{2}, \beta = 0, \gamma = 1, \delta = 0$

C)  $\alpha = 1, \beta = \frac{1}{2}, \gamma = 1, \delta = 0$

D)  $\alpha = \frac{1}{2}, \beta = \frac{1}{2}, \gamma = 1, \delta = 0$

E)  $\alpha = 1, \beta = 0, \gamma = \frac{1}{2}, \delta = 0 \checkmark$

F)  $\alpha = \frac{1}{2}, \beta = 0, \gamma = \frac{1}{2}, \delta = 0 \checkmark$

G)  $\alpha = 1, \beta = \frac{1}{2}, \gamma = \frac{1}{2}, \delta = 0 \checkmark$

H)  $\alpha = \frac{1}{2}, \beta = \frac{1}{2}, \gamma = \frac{1}{2}, \delta = 0$

I)  $\alpha = 1, \beta = 0, \gamma = 1, \delta = \frac{1}{2}$

J)  $\alpha = \frac{1}{2}, \beta = 0, \gamma = 1, \delta = \frac{1}{2}$

K)  $\alpha = 1, \beta = \frac{1}{2}, \gamma = 1, \delta = \frac{1}{2}$

L)  $\alpha = \frac{1}{2}, \beta = \frac{1}{2}, \gamma = 1, \delta = \frac{1}{2}$

M)  $\alpha = 1, \beta = 0, \gamma = \frac{1}{2}, \delta = \frac{1}{2}$

N)  $\alpha = \frac{1}{2}, \beta = 0, \gamma = \frac{1}{2}, \delta = \frac{1}{2}$

O)  $\alpha = 1, \beta = \frac{1}{2}, \gamma = \frac{1}{2}, \delta = \frac{1}{2}$

P)  $\alpha = \frac{1}{2}, \beta = \frac{1}{2}, \gamma = \frac{1}{2}, \delta = \frac{1}{2}$



15 minute(s) break!



# Playing a game

Back to the lab ☺



- We want to play a game between two solvers like this:

```
1 initial_state = State()                      # the game played
2 us = GameSolver(initial_state)                # player's solver
3 them = AnotherGameSolver(initial_state)        # opponent's solver
4 while not us.terminal():                      # until game finished
5     us.solve(), them.solve()                  # both players think
6     action = us.choose() if us.player() == 1 \
7         else them.choose()                   # one player chooses
8     us.play(action), them.play(action)        # both players act
9 out = us.outcome()                           # who won the game
10 print("won" if out > 1 else "lost" if out < -1 else "draw")
```

# Game interface (1)



- Our GameSolver's need to follow the Game interface

```
1 class Game:  
2     def reset(self):  
3         pass          # reverts game back to initial state  
4  
5     def play(self, action):  
6         pass          # changes state of the current game  
7  
8     def choose(self):  
9         return None    # ∈ A(s); chooses an action  
10  
11    def terminal(self):  
12        return True    # ∈ {T, ⊥}; whether game is finished  
13  
14    def outcome(self):  
15        return 0        # ∈ Iℝ; a finished game's outcome  
16  
17    def player(self):  
18        return 1        # ∈ {-1, 1}; which player's move it is
```

# Game implementation (1)



- The following DecisionTree can be played as a Game

```
1 class GameTree (DecisionTree, Game):
2     def __init__(self, initial_state):
3         super().__init__(initial_state)
4         self.initial_state = initial_state
5
6     def reset(self):
7         self.root = Node(self.initial_state)
8
9     def play(self, action):
10        self.root = self.get_child(self.root, action)
11
12    def get_child(self, node, action):
13        child = node.children.get(action)
14        if child is None:
15            child = Node(node.state.transition(action))
16        return child
```

## Game implementation (2)



```
1  def choose(self):
2      selected = self.best_child(self.root)
3      for action, child in self.root.children.items():
4          if child is selected:
5              return action
6
7  def best_child(self, node, **kwargs):
8      best_nodes, best_value = [], float("-inf")
9      for child in node.children.values():
10         value = self.node_value(child, parent=node, **kwargs)
11         if value == best_value:
12             best_nodes.append(child)
13         if value > best_value:
14             best_nodes, best_value = [child], value
15     return random.choice(best_nodes)
16
17 def terminal(self):
18     return self.root.state.terminal()
19
20 def outcome(self):
21     return self.root.state.reward()
```

# Game solvers



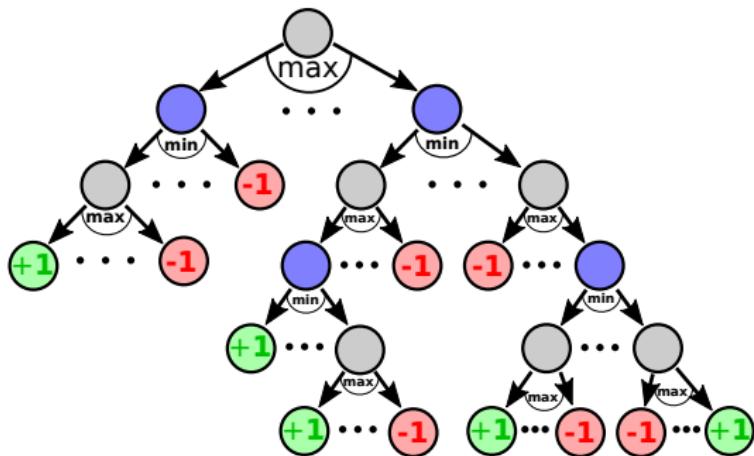
- A solver populates the DecisionTree by calling solve()

```
1 class Solver:  
2     def __init__(self, iteration_limit, **kwargs):  
3         self.iteration_limit = iteration_limit  
4         self.iteration_counter = 0  
5  
6     def solve(self):          # populate the decision tree  
7         while self.within_budget():  
8             self.search()  
9  
10    def within_budget(self):    # checks the available budget  
11        self.iteration_counter += 1  
12        done = self.iteration_counter > self.iteration_limit  
13        if done:  
14            self.iteration_counter = 0  
15        return not done  
16  
17    def search(self):          # solvers override this  
18        raise Exception("Not implemented yet.")
```

# Random search



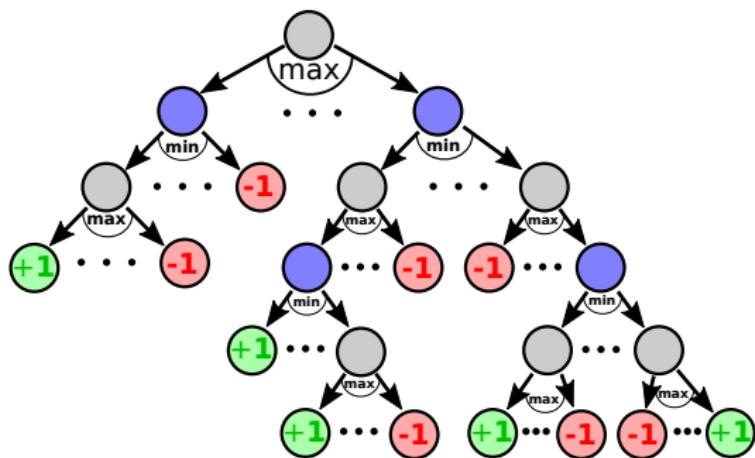
- Computing  $V(s)$  recursively can take too long



# Random search



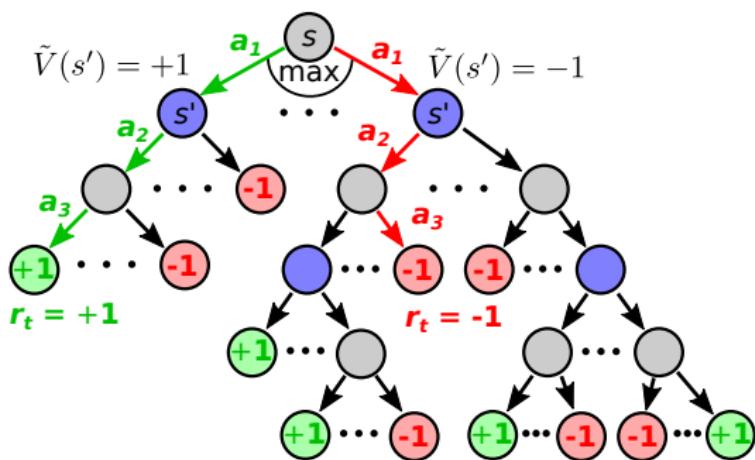
- Computing  $V(s)$  recursively can take too long
- Minimax search (plus pruning, etc) is deterministic



# Random search



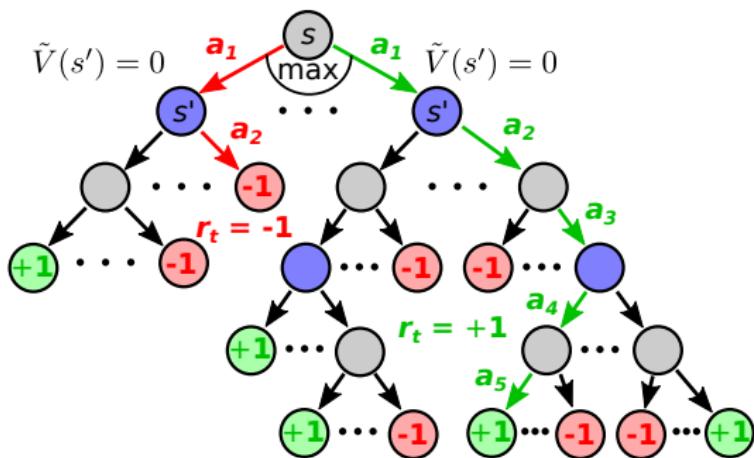
- Computing  $V(s)$  recursively can take too long
- Idea: Approximate  $V(\mathcal{P}(s, a))$ ,  $\forall a \in \mathcal{A}(s)$ , with *random* rollouts
  - keep track of  $\tilde{V}_R(s') = \frac{1}{N(s')} \sum_{t=1}^{N(s')} r_t, \forall s' \in \{\mathcal{P}(s, a) | a \in \mathcal{A}(s)\}$



# Random search



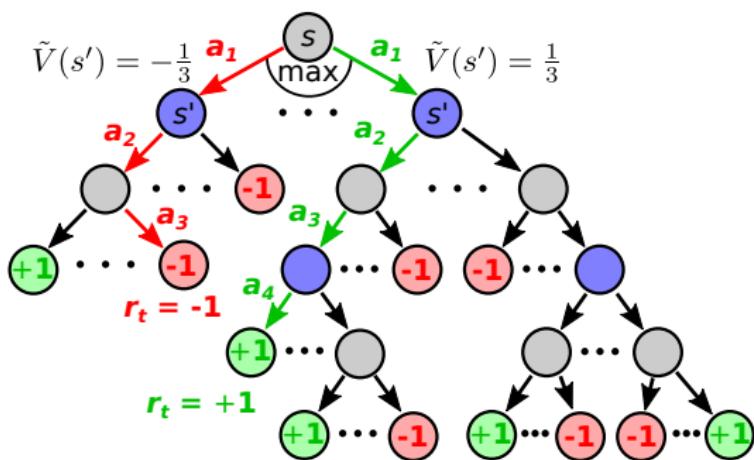
- Computing  $V(s)$  recursively can take too long
- Idea: Approximate  $V(\mathcal{P}(s, a))$ ,  $\forall a \in \mathcal{A}(s)$ , with *random* rollouts
  - keep track of  $\tilde{V}_R(s') = \frac{1}{N(s')} \sum_{t=1}^{N(s')} r_t, \forall s' \in \{\mathcal{P}(s, a) | a \in \mathcal{A}(s)\}$



# Random search



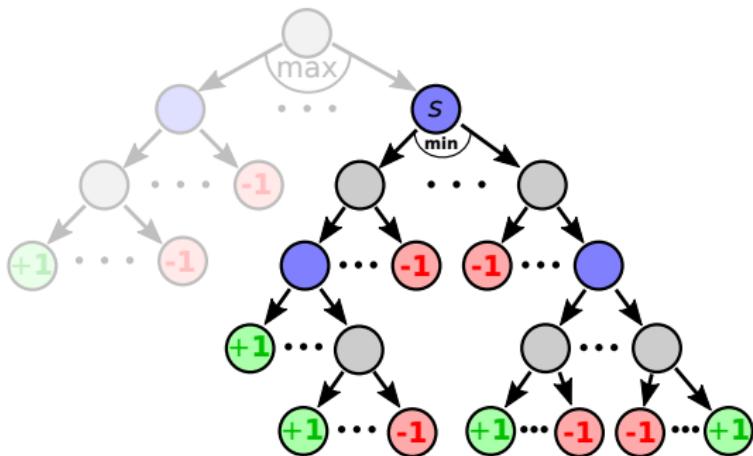
- Computing  $V(s)$  recursively can take too long
- Idea: Approximate  $V(\mathcal{P}(s, a))$ ,  $\forall a \in \mathcal{A}(s)$ , with *random* rollouts
  - keep track of  $\tilde{V}_R(s') = \frac{1}{N(s')} \sum_{t=1}^{N(s')} r_t, \forall s' \in \{\mathcal{P}(s, a) | a \in \mathcal{A}(s)\}$
  - good for shallow sub-trees with similar outcomes



# Random search



- Computing  $V(s)$  recursively can take too long
- Idea: Approximate  $V(\mathcal{P}(s, a))$ ,  $\forall a \in \mathcal{A}(s)$ , with *random* rollouts
  - keep track of  $\tilde{V}_R(s') = \frac{1}{N(s')} \sum_{t=1}^{N(s')} r_t, \forall s' \in \{\mathcal{P}(s, a) | a \in \mathcal{A}(s)\}$
  - good for shallow sub-trees with similar outcomes
  - minmax selection:  $a_t := \arg \max_{a \in \mathcal{A}(s_t)} p(s_t) \tilde{V}_R(\mathcal{P}(s_t, a))$



# Random search



- Computing  $V(s)$  recursively can take too long
- Idea: Approximate  $V(\mathcal{P}(s, a))$ ,  $\forall a \in \mathcal{A}(s)$ , with *random* rollouts
  - keep track of  $\tilde{V}_R(s') = \frac{1}{N(s')} \sum_{t=1}^{N(s')} r_t$ ,  $\forall s' \in \{\mathcal{P}(s, a) | a \in \mathcal{A}(s)\}$
  - good for shallow sub-trees with similar outcomes
  - minmax selection:  $a_t := \arg \max_{a \in \mathcal{A}(s_t)} p(s_t) \tilde{V}_R(\mathcal{P}(s_t, a))$
- Often works in practice, but is *not* Minmax Value:

$$\lim_{N(s) \rightarrow \infty} \tilde{V}_R(s) = \begin{cases} \mathcal{R}(s) & , \text{if } \mathcal{T}(s) \\ \frac{1}{\mathcal{A}(s)} \sum_{a \in \mathcal{A}(s)} \tilde{V}_R(\mathcal{P}(s, a)) & , \text{otherwise} \end{cases}$$

# Random search implementation



```
1 class RandomSearch(GameTree, Solver):
2     def search(self):                      # rollout random action
3         all_actions = self.root.state.actions()
4         action = random.choice(all_actions)
5         child = self.add_child(self.root, action)
6         child.total_value += self.rollout(child)
7         child.num_visits += 1
8
9     def add_child(self, node, action):    # potentially new node
10        child = self.get_child(node, action)
11        node.children[action] = child
12        return child
13
14    def rollout(self, node):            # evaluate random policy
15        state = node.state
16        while not state.terminal():
17            action = random.choice(state.actions())
18            state = state.transition(action)
19        return state.reward()
```



## Interactive block III: random search



Multiple (or no) answers may be correct!

### How can we exploit the tree-structure better?

- By building a decision tree using...
  - A) random rollouts from the root
  - B) random rollouts from the root's children
  - C) minmax rollouts within the tree
  - D) minmax rollouts at the tree's leafs
  - E) random rollouts at the tree's leafs
- The following nodes should be updated after a rollout:
  - F) only the root
  - G) only the root's children
  - H) all nodes along the chosen path
  - I) all nodes along the current best path
  - J) all nodes along the chosen path that differ from the current best path



## Interactive block III: random search



Multiple (or no) answers may be correct!

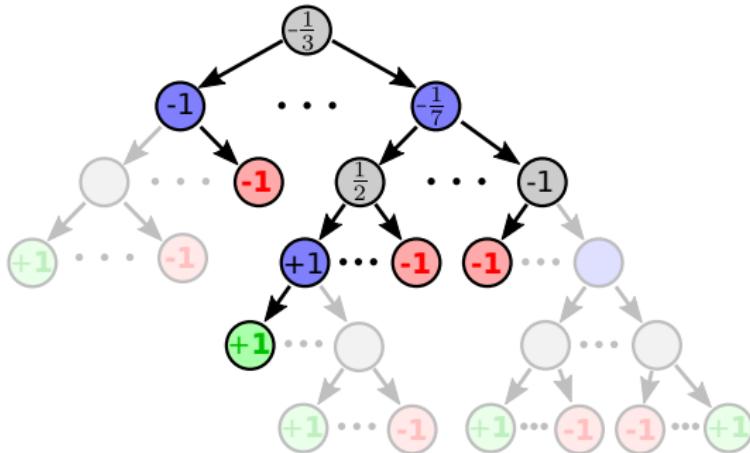
### How can we exploit the tree-structure better?

- By building a decision tree using...
  - (A) random rollouts from the root
  - (B) random rollouts from the root's children
  - (C) minmax rollouts within the tree ✓
  - (D) minmax rollouts at the tree's leafs
  - (E) random rollouts at the tree's leafs ✓
- The following nodes should be updated after a rollout:
  - (F) only the root
  - (G) only the root's children
  - (H) all nodes along the chosen path ✓
  - (I) all nodes along the current best path
  - (J) all nodes along the chosen path that differ from the current best path

# Monte Carlo Tree Search (MCTS)



- Idea: build partial tree  $\mathcal{B} \subseteq \mathcal{S}$

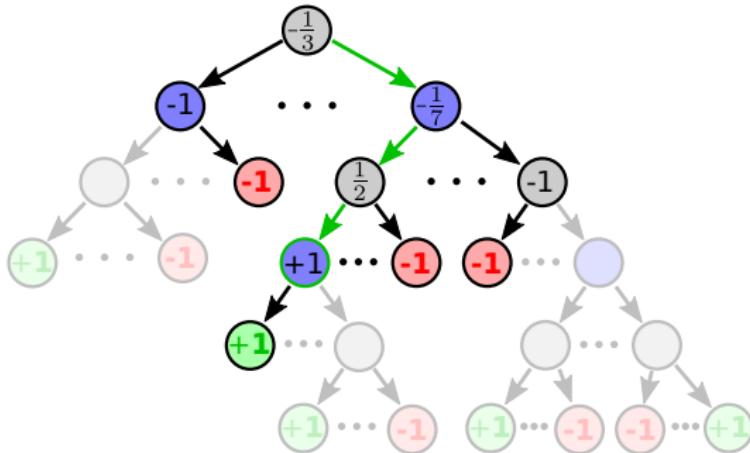


for details on MCTS see Browne et al. (2012)

# Monte Carlo Tree Search (MCTS)



- Idea: build partial tree  $\mathcal{B} \subseteq \mathcal{S}$ 
  - use minmax-selection in partial tree

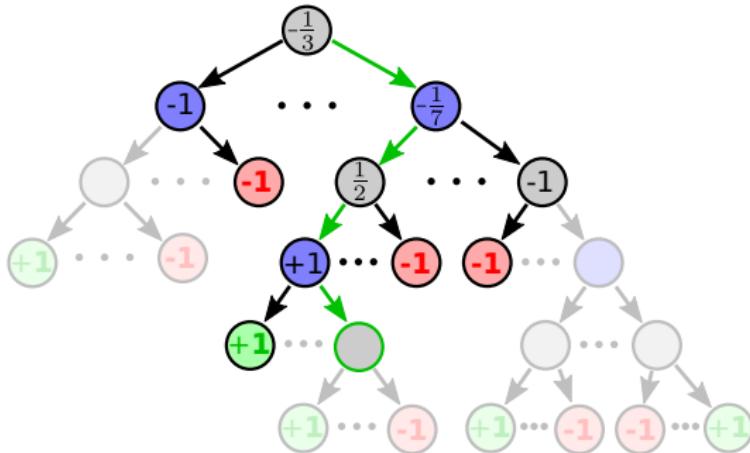


for details on MCTS see Browne et al. (2012)

# Monte Carlo Tree Search (MCTS)



- Idea: build partial tree  $\mathcal{B} \subseteq \mathcal{S}$ 
  - use minmax-selection in partial tree
  - expand tree with new leaf

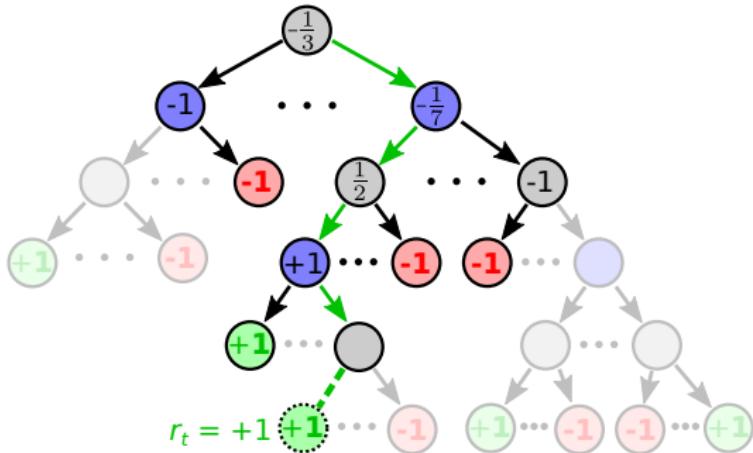


for details on MCTS see Browne et al. (2012)

# Monte Carlo Tree Search (MCTS)



- Idea: build partial tree  $\mathcal{B} \subseteq \mathcal{S}$ 
  - use minmax-selection in partial tree
  - expand tree with new leaf
  - approximate leaf values with random rollouts

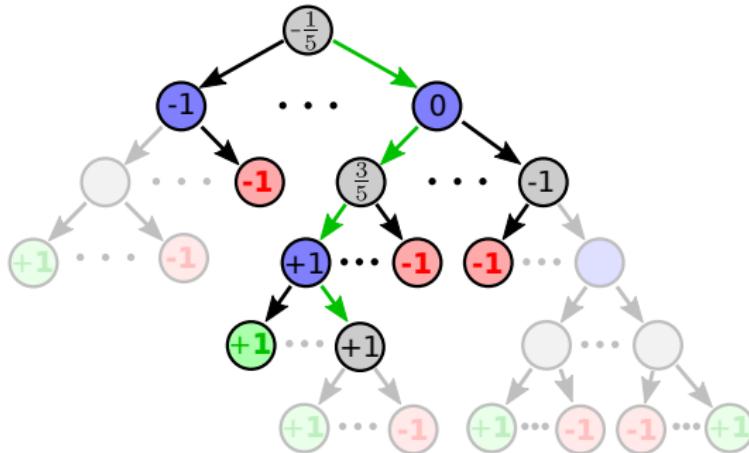


for details on MCTS see Browne et al. (2012)

# Monte Carlo Tree Search (MCTS)



- Idea: build partial tree  $\mathcal{B} \subseteq \mathcal{S}$ 
  - use minmax-selection in partial tree
  - expand tree with new leaf
  - approximate leaf values with random rollouts
  - backup the resulting reward to all parents



for details on MCTS see Browne et al. (2012)



# Monte Carlo Tree Search (MCTS)

- Idea: build partial tree  $\mathcal{B} \subseteq \mathcal{S}$ 
  - use minmax-selection in partial tree
  - expand tree with new leaf
  - approximate leaf values with random rollouts
  - backup the resulting reward to all parents
- $\mu(a|s)$  fraction of action  $a$  chosen in state  $s$ 
  - mostly minmax, but sensitive to initial rewards

$$\tilde{V}_M(s) = \begin{cases} \mathcal{R}(s) & , \text{if } \mathcal{T}(s) \\ \sum_{a \in \mathcal{A}(s)} \mu(a|s) \tilde{V}_M(\mathcal{P}(s, a)) & , \text{if } s \in \mathcal{B} \\ r_t & , \text{otherwise} \end{cases}$$

for details on MCTS see Browne et al. (2012)

# MCTS implementation (1)



- uses a stack (`list`) to keep track of visited nodes
  - ① select minmax policy
  - ② expand adds a new leaf
  - ③ rollout evaluates leaf
  - ④ backup updates values

```
1 class MCTS (RandomSearch):  
2     def __init__(self, initial_state, iteration_limit, **kwargs):  
3         super().__init__(initial_state, iteration_limit, **kwargs)  
4         self.stack = []  
5  
6     def init_stack(self, node):          # initialize stack  
7         self.stack = [node]  
8  
9     def search(self):  
10        leaf = self.select(self.root)    # push path on stack  
11        if not leaf.state.terminal():   # unless this is a leaf  
12            leaf = self.expand(leaf)    # create new leaf  
13            reward = self.rollout(leaf) # evaluate leaf  
14            self.backup(reward)       # pop path from stack
```

# MCTS implementation (2)



```
1  def select(self, node):                      # selects leaf node
2      self.init_stack(node)
3      while len(node.children) == len(node.state.actions()) \
4          and not node.state.terminal():
5          node = self.best_child(node)
6          self.stack.append(node)
7      return node
8
9  def expand(self, node):                      # adds unexplored leaf
10     for action in node.state.actions():
11         if action not in node.children:
12             leaf = self.add_child(node, action)
13             self.stack.append(leaf)
14             return leaf
15
16 def backup(self, reward):
17     while len(self.stack) > 0:
18         node = self.stack.pop()
19         node.total_value += reward
20         node.num_visits += 1
```

```
1 def search(self):
2     leaf = self.select(self.root)
3     if not leaf.state.terminal():
4         leaf = self.expand(leaf)
5         reward = self.rollout(leaf)
6         self.backup(reward)
```



# MCTS characteristics

- In MCTS  $\mu$  depends on initial rewards
  - there are ways to do proper *exploration*

$$\tilde{V}_M(s) = \begin{cases} \mathcal{R}(s) & , \text{if } \mathcal{T}(s) \\ \sum_{a \in \mathcal{A}(s)} \mu(a|s) \tilde{V}_M(\mathcal{P}(s, a)) & , \text{if } s \in \mathcal{B} \\ r_t & , \text{otherwise} \end{cases}$$

- Important characteristics of MCTS
  - aheuristic: no game-heuristics necessary
  - anytime: `within_budget()` can use time limit
  - asymmetric: explores promising sub-trees first

for more details on MCTS see Browne et al. (2012)

# MCTS characteristics

Excellent for game playing



- In MCTS  $\mu$  depends on initial rewards
  - there are ways to do proper *exploration*

$$\tilde{V}_M(s) = \begin{cases} \mathcal{R}(s) & , \text{if } \mathcal{T}(s) \\ \sum_{a \in \mathcal{A}(s)} \mu(a|s) \tilde{V}_M(\mathcal{P}(s, a)) & , \text{if } s \in \mathcal{B} \\ r_t & , \text{otherwise} \end{cases}$$

- Important characteristics of MCTS ☺
  - aheuristic: no game-heuristics necessary
  - anytime: `within_budget()` can use time limit
  - asymmetric: explores promising sub-trees first

for more details on MCTS see Browne et al. (2012)

# Example: AlphaGo

Famous in 2016



- Go is unsolvable with search alone



Image source: [wikimedia.org](#)

# Example: AlphaGo

Famous in 2016



- Go is unsolvable with search alone
  - $\mathcal{A}(s) \leq 361$ , decision tree with  $\approx 10^{768}$  nodes
  - long time horizons between action and payoff
  - more a game of pattern recognition than logic



Image source: [wikimedia.org](#)

# Example: AlphaGo

Famous in 2016



- Go is unsolvable with search alone
- AlphaGo combines search with deep neural networks:

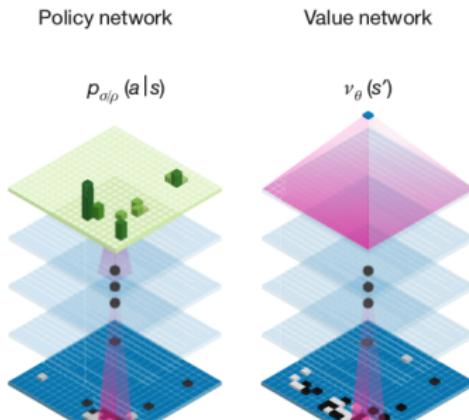
images from Silver et al. (2016)

# Example: AlphaGo

Famous in 2016



- Go is unsolvable with search alone
- AlphaGo combines search with deep neural networks:
  - MCTS with bounds (UCT)
  - neural networks as heuristics
  - behaviour cloning from human tournaments
  - further refinement through self-play



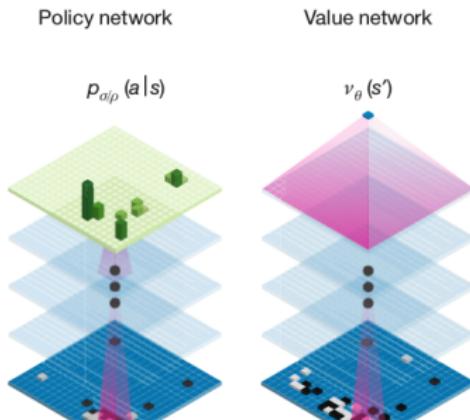
images from Silver et al. (2016)

# Example: AlphaGo

Famous in 2016



- Go is unsolvable with search alone
- AlphaGo combines search with deep neural networks:
  - MCTS with bounds (UCT)
  - neural networks as heuristics
  - ~~behaviour cloning from human tournaments~~
  - further refinement through self-play



images from Silver et al. (2016)



## Part II: Search and inference

- Today we saw monte carlo tree search



- Today we saw monte carlo tree search
  - With enhancements and NN evaluation = state of the art for go



- Today we saw monte carlo tree search
  - With enhancements and NN evaluation = state of the art for go
- Part II: constraints, propagation, search



- Today we saw monte carlo tree search
  - With enhancements and NN evaluation = state of the art for go
- Part II: constraints, propagation, search
  - state of the art for sudoku



- Today we saw monte carlo tree search
  - With enhancements and NN evaluation = state of the art for go
- Part II: constraints, propagation, search
  - state of the art for sudoku
- Lab on Monday



## Part II: Search and inference

- Today we saw monte carlo tree search
  - With enhancements and NN evaluation = state of the art for go
- Part II: constraints, propagation, search
  - state of the art for sudoku
- Lab on Monday
- Assignment 1 due: see Brightspace



- Today we saw monte carlo tree search
  - With enhancements and NN evaluation = state of the art for go
- Part II: constraints, propagation, search
  - state of the art for sudoku
- Lab on Monday
- Assignment 1 due: see Brightspace
- Questions? TUD Answers



# What is still unclear?

After every lecture...

Give us some homework:  
tell us what is still unclear.



# References I

- Cameron Browne, Edward Jack Powley, Daniel Whitehouse, Simon M. Lucas, Peter I. Cowling, Philipp Rohlfshagen, Stephen Tavener, Diego Perez Liebana, Spyridon Samothrakis, and Simon Colton. A survey of Monte Carlo tree search methods. *IEEE Trans. Comput. Intellig. and AI in Games*, 4(1):1–43, 2012.
- David Silver, Aja Huang, Chris J. Maddison, Arthur Guez, Laurent Sifre, George van den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, Sander Dieleman, Dominik Grewe, John Nham, Nal Kalchbrenner, Ilya Sutskever, Timothy Lillicrap, Madeleine Leach, Koray Kavukcuoglu, Thore Graepel, and Demis Hassabis. Mastering the game of Go with deep neural networks and tree search. *Nature*, 529(7587):484–489, January 2016.